

book

Table of contents

Streszczenie	3
Summary	3
1 Wstęp	4
I Wyniki i Dyskusja	5
2 Load data to data frames	7
2.1 Users	7
2.2 Tags	8
2.3 Votes	8
2.4 Posts	9
2.5 Post links	10
2.6 Post History	11
2.7 Badges	12
2.8 Clean up text columns	13
2.9 Write parquet	14
3 Questions/Answers over time	17
3.1 Forum retention - time from account creation to last comments	19
3.2 porównanie najwyżej i najniżej ocenianych pytań (długość, tagi, liczba odpowiedzi) . . .	23
3.3 procent przypadków kiedy najwyżej oceniana odpowiedź to nie zaakceptowana odpowiedź	33
3.4 rozkład ocen odpowiedzi zaakceptowanych vs pozostałych (średnia, odchylenie, mini- mum, maksimum)	34
3.5 top N tagów które wygenerowały najwięcej wyświetleń	35
3.6 liczba postów w czasie dla każdego z top N tagów (lineplot/barplot)	40
3.7 najczęściej pojawiające się słowa w tytułach (z pominięciem stopwords)	50
3.8 procent użytkowników którzy nigdy nic nie zapostowali	52
3.9 średni czas od pojawienia się pytania do pojawienia się zaakceptowanej odpowiedzi . . .	53
3.9.1 remove outliers	54
References	56

Streszczenie

Słowa kluczowe: Big Data, Spark, AWS, EMR, S3

Summary

Keywords: Big Data, Spark, AWS, EMR, S3

1 Wstęp

Celem niniejszej pracy jest utworzenie infrastruktury w chmurze obliczeniowej AWS pozwalające na wielkoskalową analizy danych w sytemie rozproszonym (ang. *Big Data*).

Do stworzenia przykładowego projektu wykorzystano dane ze strony *Stack Exchange* zawierającej zestawy danych pochodzące z forów społecznościowych. Analizę ograniczono do danych pochodzących z forum o nazwie *Beer, Wine and Spirits*.

W niniejszej pracy ...

Part I

Wyniki i Dyskusja

Lorem ipsum

2 Load data to data frames

```
from pyspark.sql import SparkSession
import os

PATH="jars/spark-xml_2.12-0.15.0.jar"

spark = SparkSession.builder.master("local[12]").appName("MyApp").config("spark.jars", PATH)
```

23/01/06 14:32:55 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

2.1 Users

```
from pyspark.sql.types import *

users_schema = StructType([
    StructField('_AboutMe', StringType(), True),
    StructField('_AccountId', IntegerType(), True),
    StructField('_CreationDate', TimestampType(), True),
    StructField("_DisplayName", StringType(), True),
    StructField("_DownVotes", IntegerType(), True),
    StructField("_Id", IntegerType(), True),
    StructField("_LastAccessDate", TimestampType()),
    StructField("_Location", StringType(), True),
    StructField("_ProfileImageUrl", StringType(), True),
    StructField("_Reputation", IntegerType(), True),
    StructField("_UpVotes", IntegerType(), True),
    StructField("_Views", IntegerType(), True),
    StructField("_WebsiteUrl", StringType(), True)
])

users = spark.read.format('xml').options(rowTag='row').schema(users_schema).load("data/Users")
users.show(5)
```

```
+-----+-----+-----+-----+-----+-----+
|_AboutMe|_AccountId|_CreationDate|_DisplayName|_DownVotes|_Id|_LastAccessDate|
+-----+-----+-----+-----+-----+-----+
```

<p>Hi, I'm not re...	-1 2014-01-21 17:45:...	Community	478	-1 2014-01-21
<p>Dev #2 who hel...	2 2014-01-21 20:21:...	Geoff Dalgas	0	1 2016-05-06
<p>Former Stack E...	109934 2014-01-21 20:22:...	Kasra Rahjerdi	0	2 2018-01-29
\n<p>Developer at...	37099 2014-01-21 20:22:...	Adam Lear	2	3 2021-06-04
<p>BY DAY...	237028 2014-01-21 20:22:...	Arie Litovsky	0	4 2016-12-23

only showing top 5 rows

2.2 Tags

```
tags_schema = StructType([
    StructField('_Count', IntegerType(), True),
    StructField('_ExcerptPostId', IntegerType(), True),
    StructField('_Id', IntegerType(), True),
    StructField("_TagName", StringType(), True),
    StructField("_WikiPostId", IntegerType(), True)
])

tags = spark.read.format('xml').options(rowTag='row').schema(tags_schema).load("data/Tags.xml")
tags.show(n=5)
```

_Count	_ExcerptPostId	_Id	_TagName	_WikiPostId
17	5062	1	hops	5061
85	7872	2	history	7871
69	4880	4	brewing	4879
37	5109	5	serving	5108
31	304	6	temperature	303

only showing top 5 rows

2.3 Votes

```
votes_schema = StructType([
    StructField('_BountyAmount', IntegerType(), True),
    StructField('_CreationDate', TimestampType(), True),
    StructField('_Id', IntegerType(), True),
    StructField("_PostId", StringType(), True),
    StructField("_UserId", IntegerType(), True),
    StructField("_VoteTypeId", IntegerType(), True)
])

votes = spark.read.format('xml').options(rowTag='row').schema(votes_schema).load("data/Votes.xml")
votes.show(n=5)
```


_BountyAmount	_CreationDate	_Id	_PostId	_UserId	_VoteTypeId
null	2014-01-21 00:00:00	1	1	null	2
null	2014-01-21 00:00:00	2	1	null	2
null	2014-01-21 00:00:00	3	4	null	2
null	2014-01-21 00:00:00	4	1	null	2
null	2014-01-21 00:00:00	5	4	null	2

only showing top 5 rows

2.4 Posts

```
posts_schema = StructType([
    StructField('_AcceptedAnswerId', IntegerType(), True),
    StructField('_AnswerCount', IntegerType(), True),
    StructField('_Body', StringType(), True),
    StructField("_ClosedDate", TimestampType(), True),
    StructField("_CommentCount", IntegerType(), True),
    StructField("_CommunityOwnedDate", TimestampType(), True),
    StructField("_ContentLicense", StringType(), True),
    StructField("_CreationDate", TimestampType(), True),
    StructField("_FavoriteCount", IntegerType(), True),
    StructField("_Id", IntegerType(), True),
    StructField("_LastActivityDate", TimestampType(), True),
    StructField("_LastEditDate", TimestampType(), True),
    StructField("_LastEditorDisplayName", StringType(), True),
    StructField("_LastEditorUserId", IntegerType(), True),
    StructField("_OwnerDisplayName", StringType(), True),
    StructField("_OwnerUserId", IntegerType(), True),
    StructField("_ParentId", IntegerType(), True),
    StructField("_PostTypeId", IntegerType(), True),
    StructField("_Score", IntegerType(), True),
    StructField("_Tags", StringType(), True),
    StructField("_Title", StringType(), True),
    StructField("_ViewCount", IntegerType(), True),
])
```

```
posts = spark.read.format('xml').options(rowTag='row').schema(posts_schema).load("data/Posts")
posts.show(n=1,vertical=True, truncate=False)
```

```
--RECORD 0-----
 _AcceptedAnswerId      | 4
 _AnswerCount           | 1
 _Body                  | <p>I was offered a beer the other day that was reportedly made with c
 _ClosedDate            | null
 _CommentCount          | 0
```

_CommunityOwnedDate	null
_ContentLicense	CC BY-SA 3.0
_CreationDate	2014-01-21 20:26:05.383
_FavoriteCount	null
_Id	1
_LastActivityDate	2014-01-21 22:04:34.977
_LastEditDate	2014-01-21 22:04:34.977
_LastEditorDisplayName	null
_LastEditorUserId	8
_OwnerDisplayName	null
_OwnerUserId	7
_ParentId	null
_PostTypeId	1
_Score	21
_Tags	<hops>
_Title	What is a citra hop, and how does it differ from other hops?
_ViewCount	2434

only showing top 1 row

2.5 Post links

```
links_schema = StructType([
    StructField("_CreationDate", TimestampType()),
    StructField("_Id", IntegerType()),
    StructField("_LinkTypeId", IntegerType()),
    StructField("_PostId", IntegerType()),
    StructField("_RelatedPostId", IntegerType())
])
```

```
links = spark.read.format('xml').options(rowTag='row').schema(links_schema).load("data/PostL
links.show(n=2,vertical=True,truncate=False)
```

-RECORD 0-----

_CreationDate	2014-01-21 21:04:25.23
_Id	25
_LinkTypeId	3
_PostId	29
_RelatedPostId	25

-RECORD 1-----

_CreationDate	2014-01-21 21:42:09.103
_Id	89
_LinkTypeId	1
_PostId	83
_RelatedPostId	50

only showing top 2 rows

2.6 Post History

```
history_schema = StructType([
    StructField("_Comment", StringType()),
    StructField("_ContentLicense", StringType()),
    StructField("_CreationDate", TimestampType()),
    StructField("_Id", IntegerType()),
    StructField("_PostHistoryTypeId", IntegerType()),
    StructField("_PostId", IntegerType()),
    StructField("_RevisionGUID", StringType()),
    StructField("_Text", StringType()),
    StructField("_UserDisplayName", StringType()),
    StructField("_UserId", IntegerType()),
])

history = spark.read.format('xml').options(rowTag='row').schema(history_schema).load("data/P
history.show(n=5,vertical=True, truncate=False)
```

```
-RECORD 0-----
 _Comment          | null
 _ContentLicense   | CC BY-SA 3.0
 _CreationDate     | 2014-01-21 20:26:05.383
 _Id               | 1
 _PostHistoryTypeId | 2
 _PostId           | 1
 _RevisionGUID     | a17002a0-00b0-417b-a404-0d8864bbbca5
 _Text             | I was offered a beer the other day that was reportedly made with citrate
 _UserDisplayName  | null
 _UserId           | 7
-RECORD 1-----
 _Comment          | null
 _ContentLicense   | CC BY-SA 3.0
 _CreationDate     | 2014-01-21 20:26:05.383
 _Id               | 2
 _PostHistoryTypeId | 1
 _PostId           | 1
 _RevisionGUID     | a17002a0-00b0-417b-a404-0d8864bbbca5
 _Text             | What is a citra hop, and how does it differ from other hops?
 _UserDisplayName  | null
 _UserId           | 7
-RECORD 2-----
 _Comment          | null
 _ContentLicense   | CC BY-SA 3.0
 _CreationDate     | 2014-01-21 20:26:05.383
 _Id               | 3
 _PostHistoryTypeId | 3
 _PostId           | 1
 _RevisionGUID     | a17002a0-00b0-417b-a404-0d8864bbbca5
```

```

_Text          | <hops>
_UserDisplayName | null
_UserId        | 7
-RECORD 3-----
_Comment       | null
_ContentLicense | CC BY-SA 3.0
_CreationDate  | 2014-01-21 20:27:29.797
_Id           | 4
_PostHistoryTypeId | 2
_PostId        | 2
_RevisionGUID  | 128709c5-8789-4d1c-a799-49a7e37da36b
_Text          | As far as we know, when did humans first brew beer, and where? Around whe
_UserDisplayName | null
_UserId        | 7
-RECORD 4-----
_Comment       | null
_ContentLicense | CC BY-SA 3.0
_CreationDate  | 2014-01-21 20:27:29.797
_Id           | 5
_PostHistoryTypeId | 1
_PostId        | 2
_RevisionGUID  | 128709c5-8789-4d1c-a799-49a7e37da36b
_Text          | When was the first beer ever brewed?
_UserDisplayName | null
_UserId        | 7
only showing top 5 rows

```

2.7 Badges

```

badges_schema = StructType([
    StructField("_Class", IntegerType()),
    StructField("_Date", TimestampType()),
    StructField("_Id", IntegerType()),
    StructField("_Name", StringType()),
    StructField("_TagBased", BooleanType()),
    StructField("_UserId", IntegerType()),
])

badges = spark.read.format('xml').options(rowTag='row').schema(badges_schema).load("data/Badges.xml")
badges.show(n=5, vertical=True, truncate=False)

```

```

-RECORD 0-----
_Class      | 3
_Date       | 2014-01-21 20:52:16.97
_Id         | 1
_Name       | Autobiographer
_TagBased   | false

```

```

  _UserId    | 1
-RECORD 1-----
  _Class     | 3
  _Date      | 2014-01-21 20:52:16.97
  _Id        | 2
  _Name      | Autobiographer
  _TagBased  | false
  _UserId    | 2
-RECORD 2-----
  _Class     | 3
  _Date      | 2014-01-21 20:52:16.97
  _Id        | 3
  _Name      | Autobiographer
  _TagBased  | false
  _UserId    | 6
-RECORD 3-----
  _Class     | 3
  _Date      | 2014-01-21 20:52:16.97
  _Id        | 4
  _Name      | Autobiographer
  _TagBased  | false
  _UserId    | 7
-RECORD 4-----
  _Class     | 3
  _Date      | 2014-01-21 20:52:16.97
  _Id        | 5
  _Name      | Autobiographer
  _TagBased  | false
  _UserId    | 9

```

only showing top 5 rows

2.8 Clean up text columns

```

from pyspark.sql.functions import regexp_replace, trim, udf, col

from bs4 import BeautifulSoup
from html import unescape

def tags_remove(s):
    if s is not None:
        soup = BeautifulSoup(unescape(s), 'lxml')
        return soup.text
    else:
        return None
udf_tags_remove = udf(lambda m: tags_remove(m))

users_clean = users.withColumn("_AboutMe_clean", regexp_replace("_AboutMe", "\n|\t|\r", " "))

```

```

.withColumn("_AboutMe_clean", udf_tags_remove(col('_AboutMe_clean'))) \
.withColumn("_AboutMe_clean", regexp_replace("_AboutMe_clean", "\s{2,}", " ")) \
.withColumn("_AboutMe_clean", trim("_AboutMe_clean"))

history_clean = history.withColumn("_Text_clean", regexp_replace("_Text", "\n|\t|\r", " ")) \
.withColumn("_Text_clean", udf_tags_remove(col('_Text_clean'))) \
.withColumn("_Text_clean", regexp_replace("_Text_clean", "\s{2,}", " ")) \
.withColumn("_Text_clean", trim("_Text_clean"))

posts_clean = posts.withColumn("_Body_clean", regexp_replace("_Body", "\n|\t|\r", " ")) \
.withColumn("_Body_clean", udf_tags_remove(col('_Body_clean'))) \
.withColumn("_Body_clean", regexp_replace("_Body_clean", "\s{2,}", " ")) \
.withColumn("_Body_clean", trim("_Body_clean"))

```

2.9 Write parquet

```

users_clean.select(
    col("_AboutMe").alias("about_me"),
    col("_AboutMe_clean").alias("about_me_clean"),
    col("_CreationDate").alias("creation_date"),
    col("_DisplayName").alias("display_name"),
    col("_DownVotes").alias("down_votes"),
    col("_Id").alias("id"),
    col("_LastAccessDate").alias("last_access_date"),
    col("_Location").alias("location"),
    col("_ProfileImageUrl").alias("profile_image_url"),
    col("_Reputation").alias("reputatio"),
    col("_UpVotes").alias("up_votes"),
    col("_Views").alias("views"),
    col("_WebsiteUrl").alias("website_url")
).write.mode('overwrite').format('parquet').option('path', "outputs/users").save()

```

/config/workspace/.venv/lib/python3.10/site-packages/bs4/__init__.py:435: MarkupResemblesLocatorWarning: warnings.warn(

```

tags.select(
    col("_Count").alias("count"),
    col("_ExcerptPostId").alias("excerpt_post_id"),
    col("_Id").alias("id"),
    col("_TagName").alias("tag_name"),
    col("_WikiPostId").alias("wiki_post_id"),
).write.mode('overwrite').format('parquet').option('path', "outputs/tags").save()

```

```

votes.select(
    col("_BountyAmount").alias("bounty_amount"),
    col("_CreationDate").alias("creation_date"),
    col("_Id").alias("id"),
    col("_PostId").alias("post_id"),
    col("_UserId").alias("user_id"),
    col("_VoteTypeId").alias("vote_type_id"),
).write.mode('overwrite').format('parquet').option('path', "outputs/votes").save()

```

```

posts_clean.select(
    col("_AcceptedAnswerId").alias("accepted_answer_id"),
    col("_AnswerCount").alias("answer_count"),
    col("_Body").alias("body"),
    col("_Body_clean").alias("body_clean"),
    col("_ClosedDate").alias("closed_date"),
    col("_CommentCount").alias("comment_count"),
    col("_CommunityOwnedDate").alias("community_owned_date"),
    col("_ContentLicense").alias("content_licence"),
    col("_CreationDate").alias("creation_date"),
    col("_FavoriteCount").alias("favourite_count"),
    col("_Id").alias("id"),
    col("_LastActivityDate").alias("last_activity_date"),
    col("_LastEditDate").alias("last_edit_date"),
    col("_LastEditorDisplayName").alias("last_editor_display_name"),
    col("_LastEditorUserId").alias("last_editor_user_id"),
    col("_OwnerUserId").alias("owner_user_id"),
    col("_PostTypeId").alias("post_type_id"),
    col("_ParentId").alias("parent_id"),
    col("_Score").alias("score"),
    col("_Tags").alias("tags"),
    col("_Title").alias("title"),
    col("_ViewCount").alias("view_count"),
).write.mode('overwrite').format('parquet').option('path', "outputs/posts").save()

```

```

links.select(
    col("_CreationDate").alias("creation_date"),
    col("_Id").alias("id"),
    col("_LinkTypeId").alias("link_type_id"),
    col("_PostId").alias("post_id"),
    col("_RelatedPostId").alias("related_post_id"),
).write.mode('overwrite').format('parquet').option('path', "outputs/post_links").save()

```

```

history_clean.select(
    col("_Comment").alias("comment"),
    col("_ContentLicense").alias("content_license"),
    col("_CreationDate").alias("creation_date"),
    col("_Id").alias("id"),
    col("_PostHistoryTypeId").alias("post_history_type_id"),
    col("_PostId").alias("post_id"),
    col("_RevisionGUID").alias("revision_guid"),
    col("_Text").alias("text"),
    col("_Text_clean").alias("text_clean"),
    col("_UserDisplayName").alias("user_display_name"),
    col("_UserId").alias("user_id"),
).write.mode('overwrite').format('parquet').option('path', "outputs/history").save()

```

```

badges.select(
    col("_Class").alias("class"),
    col("_Date").alias("date"),
    col("_Id").alias("id"),
    col("_Name").alias("name"),
    col("_TagBased").alias("tag_based"),
    col("_UserId").alias("user_id"),
).write.mode('overwrite').format('parquet').option('path', "outputs/badges").save()

```


3 Questions/Answers over time

```
from pyspark.sql import (
    SparkSession,
    functions as f
)
import matplotlib
```

```
spark = SparkSession.builder.master("local[12]").appName("Analytics").getOrCreate()
```

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

23/01/06 14:33:37 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform

23/01/06 14:33:37 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 40

```
posts = spark.read.format('parquet').load("outputs/posts")
posts.show(1, vertical=True)
```

```
-RECORD 0-----
accepted_answer_id      | 4
answer_count            | 1
body                    | <p>I was offered ...
body_clean              | I was offered a b...
closed_date             | null
comment_count           | 0
community_owned_date    | null
content_licence         | CC BY-SA 3.0
creation_date           | 2014-01-21 20:26:...
favourite_count         | null
id                      | 1
last_activity_date      | 2014-01-21 22:04:...
last_edit_date          | 2014-01-21 22:04:...
last_editor_display_name | null
last_editor_user_id     | 8
owner_user_id           | 7
post_type_id            | 1
parent_id               | null
score                   | 21
```

```
tags | <hops>
title | What is a citra h...
view_count | 2434
only showing top 1 row
```

```
posts_grouped = (
    posts
    .filter(f.col('owner_user_id').isNotNull())
    .groupBy(
        f.window('creation_date', '4 weeks')
    )
    .agg(
        f.sum(f.lit(1)).alias('all'),
        f.sum(f.when(f.col('post_type_id') == 1, f.lit(1)).otherwise(f.lit(0))).alias('questions'),
        f.sum(f.when(f.col('post_type_id') == 2, f.lit(1)).otherwise(f.lit(0))).alias('answers')
    )
    # window struct has nested columns 'start' and 'end'
    .withColumn('date', f.col('window.start').cast('date'))
    .orderBy('date')
).toPandas()
```

```
posts_grouped.head()
```

	window	all	questions	answers	date
0	(2014-01-02 00:00:00, 2014-01-30 00:00:00)	413	150	243	2014-01-02
1	(2014-01-30 00:00:00, 2014-02-27 00:00:00)	190	58	118	2014-01-30
2	(2014-02-27 00:00:00, 2014-03-27 00:00:00)	50	16	34	2014-02-27
3	(2014-03-27 00:00:00, 2014-04-24 00:00:00)	47	16	31	2014-03-27
4	(2014-04-24 00:00:00, 2014-05-22 00:00:00)	44	10	34	2014-04-24

```
# posts_grouped.plot(
#     x='date',
#     figsize=(12, 6),
#     title='Number of questions/answers per month (4 weeks)',
#     legend=True,
#     xlabel='Date',
#     ylabel='Count',
#     kind='line'
# )
```


```
from plotnine import aes, facet_wrap, ggplot, scale_x_datetime, options, stat_smooth, geom_c

options.figure_size = (15, 10)

posts_long = posts_grouped.melt(id_vars=('date'), value_vars=('all', 'questions', 'answers'))
```

```
posts_long.head()

(ggplot(posts_long, aes(x='date', y='value', group='variable'))
+ geom_col(aes(fill='variable'))
+ scale_x_datetime()
+ stat_smooth(method='loess')
+ facet_wrap('variable', ncol=1)
)
```



```
<ggplot: (8788967743071)>
```

3.1 Forum retention - time from account creation to last comments

```
users = spark.read.format('parquet').load("outputs/users").select(f.col('id'), f.col('creation_date'))

users.show()
```

```
+---+-----+-----+
| id|      creation_date|  display_name|
+---+-----+-----+
| -1|2014-01-21 17:45:...|    Community|
|  1|2014-01-21 20:21:...|    Geoff Dalgas|
```

	2	2014-01-21 20:22:...		Kasra Rahjerdi
	3	2014-01-21 20:22:...		Adam Lear
	4	2014-01-21 20:22:...		Arie Litovsky
	5	2014-01-21 20:22:...		Brian Nickel
	6	2014-01-21 20:23:...		Jeremy T
	7	2014-01-21 20:24:...		Tom Medley
	8	2014-01-21 20:25:...		LessPop_MoreFizz
	9	2014-01-21 20:25:...		Nick Craver
	10	2014-01-21 20:28:...		ChrisG
	11	2014-01-21 20:28:...		hairboat
	12	2014-01-21 20:29:...		nhaarman
	13	2014-01-21 20:29:...		Shog9
	14	2014-01-21 20:32:...		Ben Collins
	15	2014-01-21 20:33:...		Ana
	16	2014-01-21 20:34:...		Grace Note
	17	2014-01-21 20:34:...		Jon Ericson
	18	2014-01-21 20:36:...		awesome
	19	2014-01-21 20:36:...		Steve Robbins

+---+-----+-----+-----+-----+-----+

only showing top 20 rows

```
posts_by_user = posts.select(f.col('owner_user_id'), f.col('last_activity_date'), f.col('id')

posts_and_users_joined = (users
    .filter(f.col('id') != -1) # remove bots
    .join(posts_by_user, users.id == posts_by_user.owner_user_id, how="left" )
    .filter(f.col('post_id').isNotNull()) # remove users that never posted
)

posts_and_users_joined.show()
```

id	creation_date	display_name	owner_user_id	last_activity_date	post_id
2	2014-01-21 20:22:...	Kasra Rahjerdi		2 2014-01-22 00:26:...	39
2	2014-01-21 20:22:...	Kasra Rahjerdi		2 2014-01-22 05:50:...	28
4	2014-01-21 20:22:...	Arie Litovsky		4 2016-03-06 04:31:...	85
5	2014-01-21 20:22:...	Brian Nickel		5 2014-02-04 23:20:...	533
5	2014-01-21 20:22:...	Brian Nickel		5 2014-01-23 16:41:...	267
5	2014-01-21 20:22:...	Brian Nickel		5 2014-01-22 18:36:...	217
5	2014-01-21 20:22:...	Brian Nickel		5 2018-08-09 15:38:...	50
5	2014-01-21 20:22:...	Brian Nickel		5 2014-01-21 20:59:...	32
5	2014-01-21 20:22:...	Brian Nickel		5 2014-01-21 20:45:...	17
7	2014-01-21 20:24:...	Tom Medley		7 2014-01-29 20:34:...	426
7	2014-01-21 20:24:...	Tom Medley		7 2014-11-19 15:11:...	82
7	2014-01-21 20:24:...	Tom Medley		7 2014-01-21 21:47:...	70
7	2014-01-21 20:24:...	Tom Medley		7 2020-08-28 07:35:...	59
7	2014-01-21 20:24:...	Tom Medley		7 2022-01-14 10:04:...	38

	7 2014-01-21 20:24:...	Tom Medley	7 2014-01-22 06:24:...	35
	7 2014-01-21 20:24:...	Tom Medley	7 2014-01-22 17:04:...	10
	7 2014-01-21 20:24:...	Tom Medley	7 2017-08-24 06:53:...	8
	7 2014-01-21 20:24:...	Tom Medley	7 2021-01-15 06:17:...	7
	7 2014-01-21 20:24:...	Tom Medley	7 2017-06-07 11:10:...	5
	7 2014-01-21 20:24:...	Tom Medley	7 2015-01-29 14:50:...	3
+	+	+	+	+

only showing top 20 rows

```
posts_and_users_joined.select(f.col('post_id')).count() == posts_and_users_joined.select(f.c
```

True

```
user_last_post = (posts_and_users_joined
    .groupBy(f.col('id'), f.col('creation_date'))
    .agg(
        f.max(f.col('last_activity_date'))
    )
)

# time from account creation to last activity
user_last_post = user_last_post.withColumn('diff',f.datediff(f.col('max(last_activity_date)',
user_last_post.show()
```

+	+	+	+	+
	id	creation_date	max(last_activity_date)	diff
+	+	+	+	+
	6696	2017-04-27 18:46:...	2017-09-18 21:40:...	144
	7212	2017-10-24 01:20:...	2017-10-24 01:20:...	0
	7311	2017-11-28 23:29:...	2017-12-23 14:32:...	25
	10039	2020-01-17 20:47:...	2020-03-25 21:19:...	68
	149	2014-01-22 16:41:...	2014-01-23 08:52:...	1
	736	2014-04-03 13:25:...	2017-01-27 14:13:...	1030
	4197	2015-06-10 19:13:...	2015-06-10 19:18:...	0
	5654	2016-07-13 09:17:...	2016-09-02 12:08:...	51
	7154	2017-10-01 06:24:...	2017-10-01 06:24:...	0
	7286	2017-11-23 11:43:...	2017-11-24 22:07:...	1
	7936	2018-07-20 09:31:...	2018-07-24 11:49:...	4
	11698	2020-11-12 20:24:...	2020-11-12 20:24:...	0
	1295	2014-09-10 17:56:...	2016-10-08 14:17:...	759
	5893	2016-09-11 03:06:...	2016-09-11 03:07:...	0
	6636	2017-04-06 13:23:...	2017-04-06 13:23:...	0
	6699	2017-04-28 07:43:...	2017-11-16 18:08:...	202
	7208	2017-10-22 23:11:...	2017-10-22 23:48:...	0
	8088	2018-09-15 08:09:...	2018-09-15 08:19:...	0
	740	2014-04-04 15:58:...	2016-06-21 13:46:...	809
	1077	2014-07-09 23:08:...	2018-10-09 20:22:...	1553

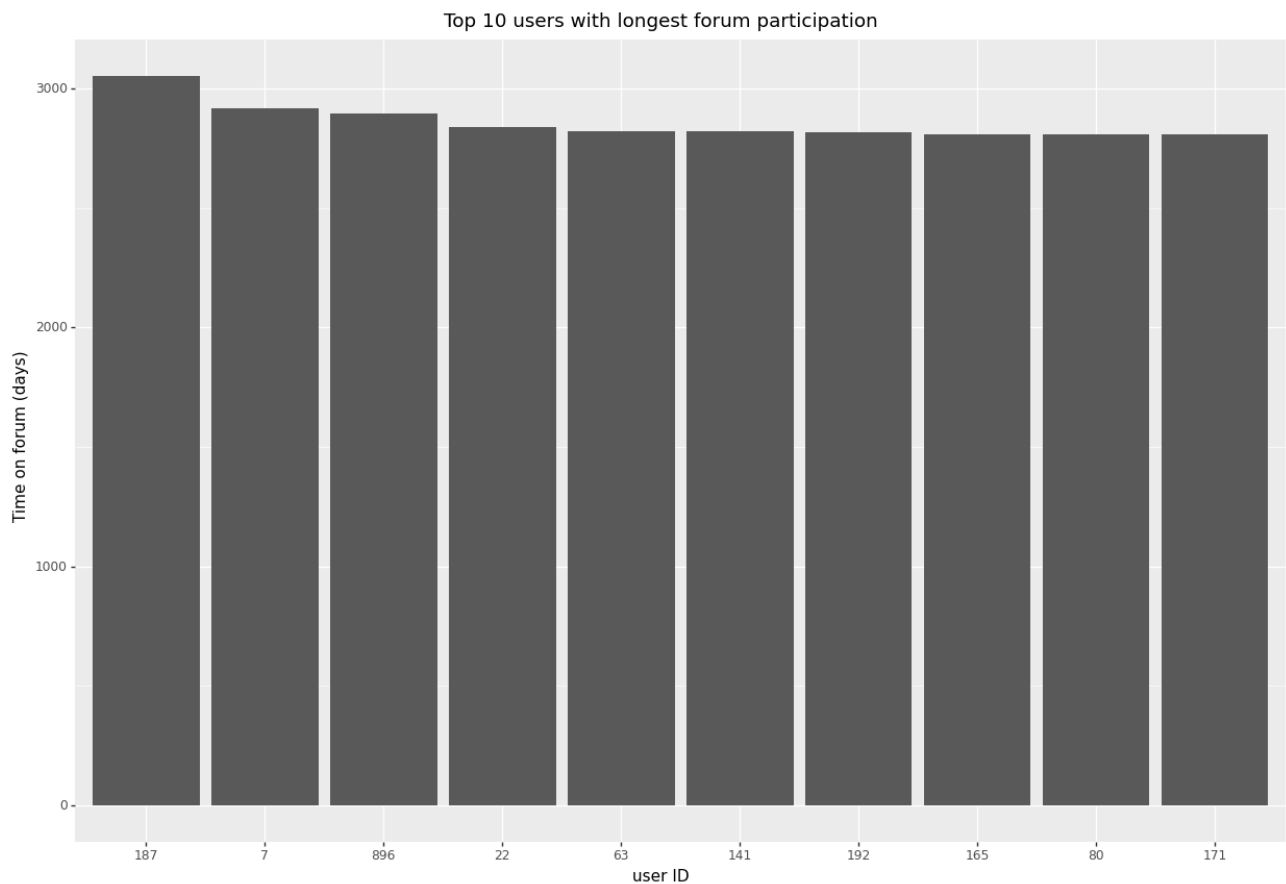
+-----+-----+-----+-----+
only showing top 20 rows

```
user_last_post.select(f.col('id')).count() == user_last_post.select(f.col('id')).distinct().
```

True

```
user_last_post_df = user_last_post.orderBy(f.col('diff').desc()).limit(10).withColumn('id_cat',  
  
import pandas as pd  
# add sorted categories for pretty plotting  
user_last_post_df['id_cat'] = pd.Categorical(user_last_post_df.id_cat, categories=user_last_  
  
from plotnine import labs  
(ggplot(user_last_post_df, aes(x='id_cat', y='diff'))  
  + geom_col()  
  + labs(x='user ID', y='Time on forum (days)', title='Top 10 users with longest forum pa  
)
```

/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: Futu
/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: Futu



```
<ggplot: (8788915768894)>
```

3.2 porównanie najwyżej i najniżej ocenianych pytań (długość, tagi, liczba odpowiedzi)

```
#users2 = spark.read.format('parquet').load("outputs/users").select(f.col('id'), f.col('crea
#posts.show(1, vertical=True)
```

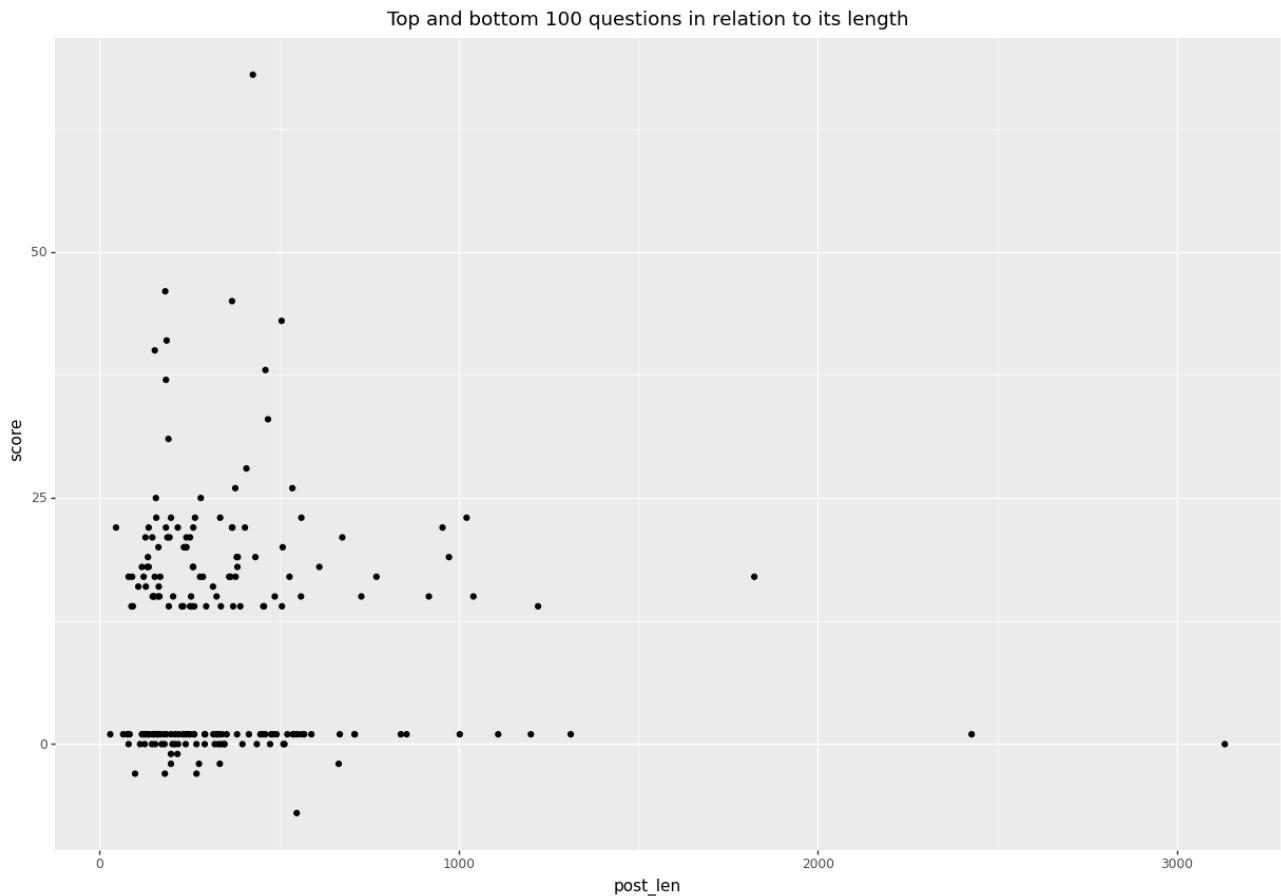
```
#1 - Question 2 - Answer 3 - Wiki 4 - TagWikiExcerpt 5 - TagWiki 6 - ModeratorNomination 7 -
questions = posts.select(f.col('id'), f.col('body_clean'), f.col('answer_count'), f.col('vie
    .filter(f.col('post_type_id') == 1)\
    .drop(f.col('post_type_id'))
```

```
n_questions = 100
top_questions = questions.orderBy(f.col('score'), ascending=False).limit(n_questions).withCo
bottom_questions = questions.orderBy(f.col('score'), ascending=True).limit(n_questions).with
edge_questions = top_questions.unionAll(bottom_questions)
```

```
from pyspark.sql.functions import length
edge_questions = edge_questions.withColumn('post_len', f.length(f.col('body_clean')))
edge_questions_pd = edge_questions.toPandas()
```

```
from plotnine import ggplot, aes, geom_point, ggtitle
```

```
(ggplot(edge_questions_pd, aes(x = 'post_len', y = 'score')) \
    + geom_point() \
    + ggtitle(f'Top and bottom {n_questions} questions in relation to its length'))
```



```
<ggplot: (8789056834491)>
```

```
edge_questions.groupby('type')\
    .agg(
        f.max(f.col('post_len')),
        f.min(f.col('post_len')),
        f.mean(f.col('post_len')),
        f.stddev(f.col('post_len')),
        f.percentile_approx(f.col('post_len'), 0.5)
    ).show()
```


```
+-----+-----+-----+-----+-----+
|  type|max(post_len)|min(post_len)|avg(post_len)|stddev_samp(post_len)|percentile_approx(post_
+-----+-----+-----+-----+-----+
|   top|          1823|          46|       349.08|    275.08858738430234|
|bottom|          3133|          30|       389.26|    424.65483595146975|
+-----+-----+-----+-----+-----+
```

```
from plotnine import ggplot, aes, geom_boxplot, ggtitle

(ggplot(edge_questions_pd, aes(x = 'type', y = 'post_len')) \
```




```
+ geom_boxplot() \
+ ggtitle(f'Top and bottom {n_questions} questions in relation to its length'))
```



<ggplot: (8788916144588)>

```
from plotnine import ggplot, aes, geom_point, ggtitle, scale_x_log10, geom_smooth

(ggplot(edge_questions_pd, aes(x = 'view_count', y = 'answer_count', fill = 'type', size = '
+ geom_point() \
+ scale_x_log10()\
+ geom_smooth()\
+ ggtitle(f'Top and bottom {n_questions} questions in relation to its views and answers'
```



<ggplot: (8788916169164)>

```
edge_questions.groupby('type')\
    .agg(
        f.max(f.col('answer_count')),
        f.min(f.col('answer_count')),
        f.mean(f.col('answer_count')),
        f.stddev(f.col('answer_count')),
        f.percentile_approx(f.col('answer_count'), 0.5)
    ).show()
```

type	max(answer_count)	min(answer_count)	avg(answer_count)	stddev_samp(answer_count)	percentile_approx(answer_count, 0.5)
top	17	1	3.96	2.581675910150324	
bottom	4	0	1.28	0.9648363026488436	

```
edge_questions.groupby('type')\
    .agg(
        f.max(f.col('view_count')),
        f.min(f.col('view_count')),
        f.mean(f.col('view_count')),
```

```

        f.stddev(f.col('view_count')),
        f.percentile_approx(f.col('view_count'), 0.5)
    ).show()

```

```

+-----+-----+-----+-----+-----+-----+
| type|max(view_count)|min(view_count)|avg(view_count)|stddev_samp(view_count)|percentile_appr|
+-----+-----+-----+-----+-----+-----+
| top|          648941|          175|        26047.65|        76276.20019090576|
|bottom|          9124|           3|         495.56|        1220.7562924910794|
+-----+-----+-----+-----+-----+-----+

```

#<https://gist.github.com/dannymeijer/be3534470b205280e52dbbcbb19a9670>

```

from pyspark.sql import DataFrame
from pyspark.sql import functions as f

```

```

def regexp_extract_all(
    df: DataFrame,
    regex: str,
    no_of_extracts: int,
    input_column_name: str,
    output_column_name: str = "output",
    empty_array_replace: bool = True,
):
    """Pyspark implementation for extracting all matches of a reg_exp_extract

```

Background

The regular implementation of `regexp_extract` (as part of `pyspark.sql.functions` module) is not capable of returning more than 1 match on a `regexp` string at a time. This function can be used to circumvent this limitation.

How it works

You can specify a ``no_of_extracts`` which will essentially run the `regexp_extract` function that number of times on the ``input_column`` of the ``df`` (``DataFrame``). In between extracts, a set of interim columns are created where every intermediate match is stored. A distinct array is created from these matches, after which the interim columns are dropped. The resulting array is stored in the defined ``output_column``. Empty strings/values in the resulting array can optionally be dropped or kept depending on how ``empty_array_replace`` is set (default is `True`).

Usage example

In the below example, we are extracting all email-addresses from a body of text.

```

The returned DataFrame will have a new ArrayType column added named `email_addresses`

> # Assuming `df` is a valid DataFrame containing a column named `text`
> email_regex = r"[\w.-]+@[ \w.-]+\.[a-zA-Z]{1,}"
> df = regexp_extract_all(df, email_regex, 6, "text", "email_addresses", True)

Parameters
-----
df: DataFrame
    Input DataFrame

regex: str
    Regexp string to extract from input DataFrame

no_of_extracts: int
    Max number of occurrences to extract

input_column_name: str
    Name of the input column

output_column_name: str
    Name of the output column (default: output)

empty_array_replace: bool
    If set to True, will replace empty arrays with null values (default: True)
"""
repeats = range(0, no_of_extracts)

# A set of interim columns are created that will be dropped afterwards
match_columns = [f"__{r}__" for r in repeats]

# Apply regexp_extract an r number of times
for r in repeats:
    df = df.withColumn(
        match_columns[r],
        f.regexp_extract(
            f.col(input_column_name),
            # the input regex string is amended with ".*?"
            # and repeated an r number of times
            # r needs to be +1 as matching groups are 1-indexed
            "".join([f"{regex}.*?" for i in range(0, r + 1)]),
            r + 1,
        ),
    )

# Create a distinct array with all empty strings removed
df = df.withColumn(
    output_column_name,
    f.array_remove(f.array_distinct(f.array(match_columns)), ""),

```

```

)

# Replace empty string with None if empty_array_replace was set
if empty_array_replace:
    df = df.withColumn(
        output_column_name,
        f.when(f.size(output_column_name) == 0, f.lit(None)).otherwise(
            f.col(output_column_name)
        ),
    )

# Drop interim columns
for c in match_columns:
    df = df.drop(c)

return df

```

```

#edge_questions.select(f.col('tags')).withColumn('tags_split', f.regexp_extract(f.col('tags'
edge_questions = regexp_extract_all(edge_questions, r'<(\w+)>', 99, "tags", "tags_split", Tr

```

```

import pyspark.rdd as rdd
h = edge_questions.filter(f.col('type') == 'top').select(f.col('tags_split')).rdd
l = edge_questions.filter(f.col('type') == 'bottom').select(f.col('tags_split')).rdd

h.flatMap(lambda x: [y if y is not None else "" for y in x])\
    .flatMap(lambda x: [x[y] for y in range(0, len(x))])\
    .map(lambda x: (x, 1))\
    .aggregateByKey(0, (lambda acc,x: acc + x ), (lambda acc1,acc2: acc1+acc2))\
    .filter(lambda x: x[1] > 1)\
    .sortBy(lambda x: x[1], ascending=False)\
    .collect()

```

```

[('taste', 17),
 ('brewing', 14),
 ('history', 12),
 ('glassware', 8),
 ('storage', 8),
 ('serving', 8),
 ('style', 8),
 ('temperature', 6),
 ('stout', 5),
 ('terminology', 5),
 ('aging', 4),
 ('health', 4),
 ('bottles', 4),


```

```
(('ingredients', 4),
 ('breweries', 3),
 ('ipa', 3),
 ('classification', 3),
 ('whiskey', 3),
 ('tripel', 3),
 ('drinking', 3),
 ('bottling', 3),
 ('flavor', 3),
 ('colour', 3),
 ('aroma', 2),
 ('freshness', 2),
 ('ale', 2),
 ('lager', 2),
 ('preservation', 2),
 ('foam', 2),
 ('dubbel', 2),
 ('skunking', 2),
 ('laws', 2),
 ('draught', 2),
 ('pouring', 2),
 ('pairing', 2),
 ('keg', 2),
 ('water', 2),
 ('trappist', 2),
 ('carbonation', 2)]
```

```
tags_rdd = h.flatMap(lambda x: [y if y is not None else "" for y in x])\
    .flatMap(lambda x: [x[y] for y in range(0, len(x))])

x = tags_rdd.collect()
tags_str = ''
for y in range(len(x)):
    tags_str += f"{x[y]} "

from wordcloud import WordCloud
import matplotlib.pyplot as plt
wc = WordCloud(background_color='white').generate(tags_str)
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wc)
plt.axis("off")
plt.tight_layout(pad = 0)
plt.show()
```



```
tags_rdd = l.flatMap(lambda x: [y if y is not None else "" for y in x])\
    .flatMap(lambda x: [x[y] for y in range(0, len(x))])

x = tags_rdd.collect()
tags_str = ''
for y in range(len(x)):
    tags_str += f"{x[y]} "

from wordcloud import WordCloud
import matplotlib.pyplot as plt
wc = WordCloud(background_color='white').generate(tags_str)
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wc)
plt.axis("off")
plt.tight_layout(pad = 0)
plt.show()
```



```
1.flatMap(lambda x: [y if y is not None else "" for y in x])\
    .flatMap(lambda x: [x[y] for y in range(0, len(x))])\
    .map(lambda x: (x, 1))\
    .aggregateByKey(0, (lambda acc,x: acc + x ), (lambda acc1,acc2: acc1+acc2))\
    .filter(lambda x: x[1] > 1)\
    .sortBy(lambda x: x[1], ascending=False)\
    .collect()
```

```
[('wine', 20),
 ('recommendations', 14),
 ('health', 11),
 ('taste', 8),
 ('history', 8),
 ('spirits', 7),
 ('breweries', 6),
 ('storage', 5),
 ('drinking', 4),
 ('flavor', 4),
 ('brewing', 4),
 ('alcohol', 4),
 ('pairing', 3),
 ('science', 3),
 ('champagne', 3),
 ('distillation', 3),
 ('temperature', 2),
 ('glassware', 2),
 ('scotch', 2),
 ('hangover', 2),
 ('water', 2),
```



```
( 'rum', 2),
( 'liquor', 2),
( 'draught', 2),
( 'vodka', 2),
( 'mead', 2),
( 'carbonation', 2),
( 'canada', 2),
( 'drink', 2),
( 'recipes', 2)]
```

3.3 procent przypadków kiedy najwyżzej oceniana odpowiedź to nie zaakceptowana odpowiedź

```
#1 - Question 2 - Answer 3 - Wiki 4 - TagWikiExcerpt 5 - TagWiki 6 - ModeratorNomination 7 -
posts_tmp = posts.select(f.col('id'), f.col("parent_id"), f.col('accepted_answer_id'), f.col('score'))
questions = posts_tmp.filter('post_type_id == 1 and answer_count > 0')\
    .select(f.col('id').alias('q_id'), f.col('accepted_answer_id'))

answers = posts_tmp.filter(f.col('post_type_id') == 2)\
    .select(f.col('id').alias('a_id'), f.col('parent_id'), f.col('score'))

from pyspark.sql import Window

window_partition_agg = Window.partitionBy("q_id")

questions.join(answers, on=questions.q_id == answers.parent_id)\
    .sort(['q_id', 'a_id'])\
    .withColumn("max_score", f.max(f.col("score")).over(window_partition_agg))\
    .filter(f.col("score") == f.col("max_score"))\
    .filter(f.col("accepted_answer_id").isNotNull())\
    .withColumn("is_accepted_best", f.col("accepted_answer_id") == f.col("a_id"))\
    .agg(
        f.sum(f.col("is_accepted_best").cast("integer")).alias("sum"),
        f.count(f.col("q_id")).alias("count")
    )\
    .withColumn("percent", (f.col("count") - f.col("sum")) / f.col("count") * 100).show()
```

```
+---+-----+-----+
|sum|count|          percent|
+---+-----+-----+
|641|  735|12.789115646258503|
+---+-----+-----+
```

3.4 rozkład ocen odpowiedzi zaakceptowanych vs pozostałych (średnia, odchylenie, minimum, maksimum)


```
window_partition_agg = Window.partitionBy("q_id")

questions.join(answers, on=questions.q_id == answers.parent_id)\
    .sort(['q_id', 'a_id'])\
    .withColumn("is_accepted", f.col("accepted_answer_id") == f.col("a_id"))\
    .groupBy(f.col("is_accepted")).agg(
        f.avg(f.col("score")).alias("avg_score"),
        f.stddev(f.col("score")).alias("std_score"),
        f.min(f.col("score")).alias("min_score"),
        f.max(f.col("score")).alias("max_score"),
        f.count("a_id")
    ).show()
```

is_accepted	avg_score	std_score	min_score	max_score	count(a_id)
null	2.7551686615886832	3.1818372333580007	-5	30	919
true	6.395043731778426	5.915949387154137	0	46	686
false	2.5841694537346713	2.7353292123298076	-4	26	897

```
accepted_df = questions.join(answers, on=questions.q_id == answers.parent_id)\
    .sort(['q_id', 'a_id'])\
    .withColumn("is_accepted", (f.col("accepted_answer_id") == f.col("a_id")).cast("string"))\
    .withColumn("is_accepted", f.when(f.col("is_accepted").isNull(), "no accepted answer").otherwise(f.col("is_accepted")))

(ggplot(accepted_df, aes(x="is_accepted", y="score"))\
    +geom_boxplot()\
    +ggtitle("Distribution of accepted vs not accepted answers (including answers for questions with no accepted answer)"))
```



<ggplot: (8788915990562)>

3.5 top N tagów które wygenerowały najwięcej wyświetleń

```
tags_views = posts.select(['tags', 'view_count'])
tags_views_agg = regexp_extract_all(tags_views, r'<(\w+)>', 99, "tags", "tags_split", True)\
    .select([f.explode(f.col('tags_split')).alias("tag"), f.col("view_count")])\
    .filter(f.col("view_count").isNotNull())\
    .groupBy('tag')\
    .agg(
        f.sum("view_count").alias("sum_views")
    )

tag_top_views = tags_views_agg.orderBy("sum_views", ascending=False).limit(20)

tag_top_views.show()
```

23/01/06 14:35:17 WARN package: Truncated the string representation of a plan since it was too
 23/01/06 14:35:17 ERROR CodeGenerator: failed to compile: org.codehaus.janino.InternalCompilerException: Compiling "GeneratedClass" in "generated.java":

```

at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:366)
at org.codehaus.janino.UnitCompiler.access$000(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$1.visitCompilationUnit(UnitCompiler.java:336)
at org.codehaus.janino.UnitCompiler$1.visitCompilationUnit(UnitCompiler.java:333)
at org.codehaus.janino.Java$CompilationUnit.accept(Java.java:363)
at org.codehaus.janino.UnitCompiler.compileUnit(UnitCompiler.java:333)
at org.codehaus.janino.SimpleCompiler.cook(SimpleCompiler.java:235)
at org.codehaus.janino.SimpleCompiler.compileToClassLoader(SimpleCompiler.java:464)
at org.codehaus.janino.ClassBodyEvaluator.compileToClass(ClassBodyEvaluator.java:314)
at org.codehaus.janino.ClassBodyEvaluator.cook(ClassBodyEvaluator.java:237)
at org.codehaus.janino.SimpleCompiler.cook(SimpleCompiler.java:205)
at org.codehaus.commons.compiler.Cookable.cook(Cookable.java:80)
at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$.org$apache$spark$sql$ca
at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$$anon$1.load(CodeGenerat
at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$$anon$1.load(CodeGenerat
at org.sparkproject.guava.cache.LocalCache$LoadingValueReference.loadFuture(LocalCache.java
at org.sparkproject.guava.cache.LocalCache$Segment.loadSync(LocalCache.java:2379)
at org.sparkproject.guava.cache.LocalCache$Segment.lockedGetOrLoad(LocalCache.java:2342)
at org.sparkproject.guava.cache.LocalCache$Segment.get(LocalCache.java:2257)
at org.sparkproject.guava.cache.LocalCache.get(LocalCache.java:4000)
at org.sparkproject.guava.cache.LocalCache.getOrLoad(LocalCache.java:4004)
at org.sparkproject.guava.cache.LocalCache$LocalLoadingCache.get(LocalCache.java:4874)
at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$.compile(CodeGenerator.s
at org.apache.spark.sql.execution.WholeStageCodegenExec.liftedTree1$1(WholeStageCodegenExec
at org.apache.spark.sql.execution.WholeStageCodegenExec.doExecute(WholeStageCodegenExec.sca
at org.apache.spark.sql.execution.SparkPlan.$anonfun$execute$1(SparkPlan.scala:194)
at org.apache.spark.sql.execution.SparkPlan.$anonfun$executeQuery$1(SparkPlan.scala:232)
at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:151)
at org.apache.spark.sql.execution.SparkPlan.executeQuery(SparkPlan.scala:229)
at org.apache.spark.sql.execution.SparkPlan.execute(SparkPlan.scala:190)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.inputRDD$lzycompute(ShuffleE
at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.inputRDD(ShuffleExchangeExe
at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.mapOutputStatisticsFuture$lz
at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.mapOutputStatisticsFuture(SH
at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.$anonfun$submitShuffleJob$1
at org.apache.spark.sql.execution.SparkPlan.$anonfun$executeQuery$1(SparkPlan.scala:232)
at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:151)
at org.apache.spark.sql.execution.SparkPlan.executeQuery(SparkPlan.scala:229)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.submitShuffleJob(ShuffleExch
at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.submitShuffleJob$(ShuffleEx
at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.submitShuffleJob(ShuffleExch
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.shuffleFuture$lzycompute(
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.shuffleFuture(QueryStageEx
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.doMaterialize(QueryStageEx
at org.apache.spark.sql.execution.adaptive.QueryStageExec.materialize(QueryStageExec.scala
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysical
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysical
at scala.collection.Iterator.foreach(Iterator.scala:943)
at scala.collection.Iterator.foreach$(Iterator.scala:943)
at scala.collection.AbstractIterator.foreach(Iterator.scala:1431)

```

```

at scala.collection.IterableLike.foreach(IterableLike.scala:74)
at scala.collection.IterableLike.foreach$(IterableLike.scala:73)
at scala.collection.AbstractIterable.foreach(Iterable.scala:56)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysicalPlan$1(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.SparkSession.withActive(SparkSession.scala:779)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.getFinalPhysicalPlan(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.withFinalPlanUpdate(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.executeCollect(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.Dataset.collectFromPlan(Dataset.scala:3868)
at org.apache.spark.sql.Dataset.$anonfun$head$1(Dataset.scala:2863)
at org.apache.spark.sql.Dataset.$anonfun$withAction$2(Dataset.scala:3858)
at org.apache.spark.sql.execution.QueryExecution$.withInternalError(QueryExecution.scala:516)
at org.apache.spark.sql.Dataset.$anonfun$withAction$1(Dataset.scala:3856)
at org.apache.spark.sql.execution.SQLExecution$. $anonfun$withNewExecutionId$6(SQLExecution.scala:164)
at org.apache.spark.sql.execution.SQLExecution$.withSQLConfPropagated(SQLExecution.scala:164)
at org.apache.spark.sql.execution.SQLExecution$. $anonfun$withNewExecutionId$1(SQLExecution.scala:164)
at org.apache.spark.sql.SparkSession.withActive(SparkSession.scala:779)
at org.apache.spark.sql.execution.SQLExecution$.withNewExecutionId(SQLExecution.scala:64)
at org.apache.spark.sql.Dataset.withAction(Dataset.scala:3856)
at org.apache.spark.sql.Dataset.head(Dataset.scala:2863)
at org.apache.spark.sql.Dataset.take(Dataset.scala:3084)
at org.apache.spark.sql.Dataset.getRows(Dataset.scala:288)
at org.apache.spark.sql.Dataset.showString(Dataset.scala:327)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:566)
at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
at py4j.Gateway.invoke(Gateway.java:282)
at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
at py4j.commands.CallCommand.execute(CallCommand.java:79)
at py4j.ClientServerConnection.waitForCommands(ClientServerConnection.java:182)
at py4j.ClientServerConnection.run(ClientServerConnection.java:106)
at java.base/java.lang.Thread.run(Thread.java:829)
Caused by: org.codehaus.janino.InternalCompilerException: Code of method "hashAgg_doAggregateWithPartitions" is not valid
at org.codehaus.janino.CodeContext.makeSpace(CodeContext.java:1051)
at org.codehaus.janino.CodeContext.write(CodeContext.java:932)
at org.codehaus.janino.UnitCompiler.writeOpcodes(UnitCompiler.java:12101)
at org.codehaus.janino.UnitCompiler.invoke(UnitCompiler.java:11878)
at org.codehaus.janino.UnitCompiler.compileGet2(UnitCompiler.java:5186)
at org.codehaus.janino.UnitCompiler.access$9100(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$16.visitMethodInvocation(UnitCompiler.java:4482)
at org.codehaus.janino.UnitCompiler$16.visitMethodInvocation(UnitCompiler.java:4455)
at org.codehaus.janino.Java$MethodInvocation.accept(Java.java:5286)
at org.codehaus.janino.UnitCompiler.compileGet(UnitCompiler.java:4455)
at org.codehaus.janino.UnitCompiler.compileGetValue(UnitCompiler.java:5683)
at org.codehaus.janino.UnitCompiler.compileBoolean2(UnitCompiler.java:4175)
at org.codehaus.janino.UnitCompiler.access$6600(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$14.visitBinaryOperation(UnitCompiler.java:4008)

```

[illegible]

```

at org.codehaus.janino.UnitCompiler.access$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1661)
at org.codehaus.janino.UnitCompiler.access$2000(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitForStatement(UnitCompiler.java:1499)
at org.codehaus.janino.UnitCompiler$6.visitForStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$ForStatement.accept(Java.java:3187)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1848)
at org.codehaus.janino.UnitCompiler.access$2200(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitWhileStatement(UnitCompiler.java:1501)
at org.codehaus.janino.UnitCompiler$6.visitWhileStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$WhileStatement.accept(Java.java:3245)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:3420)
at org.codehaus.janino.UnitCompiler.compileDeclaredMethods(UnitCompiler.java:1362)
at org.codehaus.janino.UnitCompiler.compileDeclaredMethods(UnitCompiler.java:1335)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:807)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:975)
at org.codehaus.janino.UnitCompiler.access$700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$2.visitMemberClassDeclaration(UnitCompiler.java:392)
at org.codehaus.janino.UnitCompiler$2.visitMemberClassDeclaration(UnitCompiler.java:384)
at org.codehaus.janino.Java$MemberClassDeclaration.accept(Java.java:1445)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:384)
at org.codehaus.janino.UnitCompiler.compileDeclaredMemberTypes(UnitCompiler.java:1312)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:833)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:410)
at org.codehaus.janino.UnitCompiler.access$400(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$2.visitPackageMemberClassDeclaration(UnitCompiler.java:392)
at org.codehaus.janino.UnitCompiler$2.visitPackageMemberClassDeclaration(UnitCompiler.java:384)
at org.codehaus.janino.Java$PackageMemberClassDeclaration.accept(Java.java:1594)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:384)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:362)
... 84 more

```

```

23/01/06 14:35:17 WARN WholeStageCodegenExec: Whole-stage codegen disabled for plan (id=1):
  *(1) HashAggregate(keys=[tag#23691], functions=[partial_sum(view_count#21)], output=[tag#23691])
+- *(1) Project [tag#23691, view_count#21]
  +- *(1) Generate explode(tags_split#18439), [view_count#21], false, [tag#23691]
    +- *(1) Project [view_count#21, CASE WHEN (size(tags_split#18336, true) = 0) THEN null ELSE ...

```

```

+- *(1) Project [view_count#21, array_remove(array_distinct(array(regexp_extract(tags#
+- *(1) Filter ((isnotnull(view_count#21) AND (size(CASE WHEN (size(array_remove(ar
+- *(1) ColumnarToRow
+- FileScan parquet [tags#19,view_count#21] Batched: true, DataFilters: [isno

```

```

+-----+-----+
|          tag|sum_views|
+-----+-----+
|          taste| 1330670|
|          health| 1286001|
|  preservation|  682216|
|          storage|  542860|
|          whiskey|  464756|
|          bourbon|  330268|
|          brewing|  307892|
|           ipa|  291935|
|          spirits|  255328|
|          drinking|  225924|
| temperature|  218203|
|          drink|  204991|
|          tequila|  196689|
|          alcohol|  188615|
| recommendations|  185154|
|           wine|  181081|
|          style|  168681|
|          flavor|  168594|
|          history|  167414|
|          pairing|  164589|
+-----+-----+

```

3.6 liczba postów w czasie dla każdego z top N tagów (lineplot/barplot)

```

#1 - Question 2 - Answer 3 - Wiki 4 - TagWikiExcerpt 5 - TagWiki 6 - ModeratorNomination 7 -
posts_tmp = posts.select(f.col('id'), f.col('creation_date'), f.col('tags'))

posts_tags_time = regexp_extract_all(posts_tmp, r'<(\w+)>', 99, "tags", "tags_split", True).

top_posts_tags_time = posts_tags_time.join(tag_top_views, on="tag", how="inner").select(f.co

from plotnine import geom_bar, ylim

top_posts_tags_time_agg = top_posts_tags_time.groupBy(
    f.window('creation_date', '4 weeks'), f.col("tag")
) \
    .agg(
        f.sum(f.lit(1)).alias('count')
    ) \

```



```

        .withColumn('date', f.col('window.start').cast('date'))

top_posts_tags_time_agg_pd = top_posts_tags_time_agg.toPandas()

(ggplot(top_posts_tags_time_agg_pd, aes("date", "count"))\
  + scale_x_datetime()\
  + geom_col() \
  + facet_wrap("tag", ncol=3) \
  + ylim(0, 15))

```

```

23/01/06 14:35:22 ERROR CodeGenerator: failed to compile: org.codehaus.janino.InternalCompilerException: Compiling "GeneratedClass" in "generated.java":
    at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:366)
    at org.codehaus.janino.UnitCompiler.access$000(UnitCompiler.java:226)
    at org.codehaus.janino.UnitCompiler$1.visitCompilationUnit(UnitCompiler.java:336)
    at org.codehaus.janino.UnitCompiler$1.visitCompilationUnit(UnitCompiler.java:333)
    at org.codehaus.janino.Java$CompilationUnit.accept(Java.java:363)
    at org.codehaus.janino.UnitCompiler.compileUnit(UnitCompiler.java:333)
    at org.codehaus.janino.SimpleCompiler.cook(SimpleCompiler.java:235)
    at org.codehaus.janino.SimpleCompiler.compileToClassLoader(SimpleCompiler.java:464)
    at org.codehaus.janino.ClassBodyEvaluator.compileToClass(ClassBodyEvaluator.java:314)
    at org.codehaus.janino.ClassBodyEvaluator.cook(ClassBodyEvaluator.java:237)
    at org.codehaus.janino.SimpleCompiler.cook(SimpleCompiler.java:205)
    at org.codehaus.commons.compiler.Cookable.cook(Cookable.java:80)
    at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$.org$apache$spark$sql$codegen$CodeGenerator$$anon$1.load(CodeGenerator.scala:114)
    at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$.org$apache$spark$sql$codegen$CodeGenerator$$anon$1.load(CodeGenerator.scala:114)
    at org.sparkproject.guava.cache.LocalCache$LoadingValueReference.loadFuture(LocalCache.java:2539)
    at org.sparkproject.guava.cache.LocalCache$Segment.loadSync(LocalCache.java:2379)
    at org.sparkproject.guava.cache.LocalCache$Segment.lockedGetOrLoad(LocalCache.java:2342)
    at org.sparkproject.guava.cache.LocalCache$Segment.get(LocalCache.java:2257)
    at org.sparkproject.guava.cache.LocalCache.get(LocalCache.java:4000)
    at org.sparkproject.guava.cache.LocalCache.getOrLoad(LocalCache.java:4004)
    at org.sparkproject.guava.cache.LocalCache$LocalLoadingCache.get(LocalCache.java:4874)
    at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$.compile(CodeGenerator.scala:114)
    at org.apache.spark.sql.execution.WholeStageCodegenExec.liftedTree1$1(WholeStageCodegenExec.scala:73)
    at org.apache.spark.sql.execution.WholeStageCodegenExec.doExecute(WholeStageCodegenExec.scala:73)
    at org.apache.spark.sql.execution.SparkPlan.$anonfun$execute$1(SparkPlan.scala:194)
    at org.apache.spark.sql.execution.SparkPlan.$anonfun$executeQuery$1(SparkPlan.scala:232)
    at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:151)
    at org.apache.spark.sql.execution.SparkPlan.executeQuery(SparkPlan.scala:229)
    at org.apache.spark.sql.execution.SparkPlan.execute(SparkPlan.scala:190)
    at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.inputRDD$lzycompute(ShuffleExchangeExec.scala:100)
    at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.inputRDD(ShuffleExchangeExec.scala:100)
    at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.mapOutputStatisticsFuture$lzycompute(ShuffleExchangeExec.scala:100)
    at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.mapOutputStatisticsFuture(ShuffleExchangeExec.scala:100)
    at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.$anonfun$submitShuffleJob$1(ShuffleExchangeLike.scala:100)
    at org.apache.spark.sql.execution.SparkPlan.$anonfun$executeQuery$1(SparkPlan.scala:232)
    at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:151)

```

```

at org.apache.spark.sql.execution.SparkPlan.executeQuery(SparkPlan.scala:229)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.submitShuffleJob(ShuffleExchangeLike.scala:100)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.submitShuffleJob$(ShuffleExchangeLike.scala:100)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.submitShuffleJob(ShuffleExchangeExec.scala:100)
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.shuffleFuture$lzycompute(ShuffleQueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.shuffleFuture(ShuffleQueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.doMaterialize(ShuffleQueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.QueryStageExec.materialize(QueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysicalPlan$1(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysicalPlan$1(AdaptiveSparkPlanExec.scala:100)
at scala.collection.Iterator.foreach(Iterator.scala:943)
at scala.collection.Iterator.foreach$(Iterator.scala:943)
at scala.collection.AbstractIterator.foreach(Iterator.scala:1431)
at scala.collection.IterableLike.foreach(IterableLike.scala:74)
at scala.collection.IterableLike.foreach$(IterableLike.scala:73)
at scala.collection.AbstractIterable.foreach(Iterable.scala:56)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysicalPlan$1(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.SparkSession.withActive(SparkSession.scala:779)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.getFinalPhysicalPlan(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.withFinalPlanUpdate(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.executeCollect(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.Dataset.$anonfun$collectToPython$1(Dataset.scala:3688)
at org.apache.spark.sql.Dataset.$anonfun$withAction$2(Dataset.scala:3858)
at org.apache.spark.sql.execution.QueryExecution$.withInternalError(QueryExecution.scala:51)
at org.apache.spark.sql.Dataset.$anonfun$withAction$1(Dataset.scala:3856)
at org.apache.spark.sql.execution.SQLExecution$. $anonfun$withNewExecutionId$6(SQLExecution.scala:16)
at org.apache.spark.sql.execution.SQLExecution$.withSQLConfPropagated(SQLExecution.scala:16)
at org.apache.spark.sql.execution.SQLExecution$. $anonfun$withNewExecutionId$1(SQLExecution.scala:16)
at org.apache.spark.sql.SparkSession.withActive(SparkSession.scala:779)
at org.apache.spark.sql.execution.SQLExecution$.withNewExecutionId(SQLExecution.scala:64)
at org.apache.spark.sql.Dataset.withAction(Dataset.scala:3856)
at org.apache.spark.sql.Dataset.collectToPython(Dataset.scala:3685)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:566)
at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
at py4j.Gateway.invoke(Gateway.java:282)
at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
at py4j.commands.CallCommand.execute(CallCommand.java:79)
at py4j.ClientServerConnection.waitForCommands(ClientServerConnection.java:182)
at py4j.ClientServerConnection.run(ClientServerConnection.java:106)
at java.base/java.lang.Thread.run(Thread.java:829)
Caused by: org.codehaus.janino.InternalCompilerException: Code of method "hashAgg_doAggregateWithPartitions" is not valid
at org.codehaus.janino.CodeContext.makeSpace(CodeContext.java:1051)
at org.codehaus.janino.CodeContext.write(CodeContext.java:932)
at org.codehaus.janino.UnitCompiler.writeOpcode(UnitCompiler.java:12101)
at org.codehaus.janino.UnitCompiler.invoke(UnitCompiler.java:11878)
at org.codehaus.janino.UnitCompiler.compileGet2(UnitCompiler.java:5186)

```

at org.codehaus.janino.UnitCompiler.access\$9100(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$16.visitMethodInvocation(UnitCompiler.java:4482)
at org.codehaus.janino.UnitCompiler\$16.visitMethodInvocation(UnitCompiler.java:4455)
at org.codehaus.janino.Java\$MethodInvocation.accept(Java.java:5286)
at org.codehaus.janino.UnitCompiler.compileGet(UnitCompiler.java:4455)
at org.codehaus.janino.UnitCompiler.compileGetValue(UnitCompiler.java:5683)
at org.codehaus.janino.UnitCompiler.compileBoolean2(UnitCompiler.java:4175)
at org.codehaus.janino.UnitCompiler.access\$6600(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$14.visitBinaryOperation(UnitCompiler.java:4008)
at org.codehaus.janino.UnitCompiler\$14.visitBinaryOperation(UnitCompiler.java:3986)
at org.codehaus.janino.Java\$BinaryOperation.accept(Java.java:5077)
at org.codehaus.janino.UnitCompiler.compileBoolean(UnitCompiler.java:3986)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:2474)
at org.codehaus.janino.UnitCompiler.access\$1900(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitIfStatement(UnitCompiler.java:1498)
at org.codehaus.janino.UnitCompiler\$6.visitIfStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$IfStatement.accept(Java.java:3140)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access\$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:2475)
at org.codehaus.janino.UnitCompiler.access\$1900(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitIfStatement(UnitCompiler.java:1498)
at org.codehaus.janino.UnitCompiler\$6.visitIfStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$IfStatement.accept(Java.java:3140)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access\$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:2486)
at org.codehaus.janino.UnitCompiler.access\$1900(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitIfStatement(UnitCompiler.java:1498)
at org.codehaus.janino.UnitCompiler\$6.visitIfStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$IfStatement.accept(Java.java:3140)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access\$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$Block.accept(Java.java:2969)

```

at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1598)
at org.codehaus.janino.UnitCompiler.access$2600(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitDoStatement(UnitCompiler.java:1505)
at org.codehaus.janino.UnitCompiler$6.visitDoStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$DoStatement.accept(Java.java:3664)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1661)
at org.codehaus.janino.UnitCompiler.access$2000(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitForStatement(UnitCompiler.java:1499)
at org.codehaus.janino.UnitCompiler$6.visitForStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$ForStatement.accept(Java.java:3187)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1848)
at org.codehaus.janino.UnitCompiler.access$2200(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitWhileStatement(UnitCompiler.java:1501)
at org.codehaus.janino.UnitCompiler$6.visitWhileStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$WhileStatement.accept(Java.java:3245)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:3420)
at org.codehaus.janino.UnitCompiler.compileDeclaredMethods(UnitCompiler.java:1362)
at org.codehaus.janino.UnitCompiler.compileDeclaredMethods(UnitCompiler.java:1335)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:807)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:975)
at org.codehaus.janino.UnitCompiler.access$700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$2.visitMemberClassDeclaration(UnitCompiler.java:392)
at org.codehaus.janino.UnitCompiler$2.visitMemberClassDeclaration(UnitCompiler.java:384)
at org.codehaus.janino.Java$MemberClassDeclaration.accept(Java.java:1445)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:384)
at org.codehaus.janino.UnitCompiler.compileDeclaredMemberTypes(UnitCompiler.java:1312)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:833)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:410)
at org.codehaus.janino.UnitCompiler.access$400(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$2.visitPackageMemberClassDeclaration(UnitCompiler.java:392)
at org.codehaus.janino.UnitCompiler$2.visitPackageMemberClassDeclaration(UnitCompiler.java:384)

```

```

    at org.codehaus.janino.Java$PackageMemberClassDeclaration.accept(Java.java:1594)
    at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:384)
    at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:362)
    ... 80 more
23/01/06 14:35:22 WARN WholeStageCodegenExec: Whole-stage codegen disabled for plan (id=1):
  *(1) HashAggregate(keys=[tag#23691], functions=[partial_sum(view_count#34545)], output=[tag#23691])
+- *(1) Project [tag#23691, view_count#34545]
  +- *(1) Generate explode(tags_split#18439), [view_count#34545], false, [tag#23691]
    +- *(1) Project [view_count#34545, CASE WHEN (size(tags_split#18336, true) = 0) THEN null]
      +- *(1) Project [view_count#34545, array_remove(array_distinct(array(regexp_extract(tags_split#18439, '([a-zA-Z0-9_]+)', 1)), view_count#34545))]
        +- *(1) Filter ((isnotnull(view_count#34545) AND (size(CASE WHEN (size(array_remove(tags_split#18439, view_count#34545), 0) = 0) THEN null, array_remove(tags_split#18439, view_count#34545)) = 0) THEN null))
          +- *(1) ColumnarToRow
            +- FileScan parquet [tags#34543,view_count#34545] Batched: true, DataFilters: []
23/01/06 14:35:23 ERROR CodeGenerator: failed to compile: org.codehaus.janino.InternalCompilerException:
org.codehaus.janino.InternalCompilerException: Compiling "GeneratedClass" in "generated.java":
  at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:366)
  at org.codehaus.janino.UnitCompiler.access$000(UnitCompiler.java:226)
  at org.codehaus.janino.UnitCompiler$1.visitCompilationUnit(UnitCompiler.java:336)
  at org.codehaus.janino.UnitCompiler$1.visitCompilationUnit(UnitCompiler.java:333)
  at org.codehaus.janino.Java$CompilationUnit.accept(Java.java:363)
  at org.codehaus.janino.UnitCompiler.compileUnit(UnitCompiler.java:333)
  at org.codehaus.janino.SimpleCompiler.cook(SimpleCompiler.java:235)
  at org.codehaus.janino.SimpleCompiler.compileToClassLoader(SimpleCompiler.java:464)
  at org.codehaus.janino.ClassBodyEvaluator.compileToClass(ClassBodyEvaluator.java:314)
  at org.codehaus.janino.ClassBodyEvaluator.cook(ClassBodyEvaluator.java:237)
  at org.codehaus.janino.SimpleCompiler.cook(SimpleCompiler.java:205)
  at org.codehaus.commons.compiler.Cookable.cook(Cookable.java:80)
  at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$.org$apache$spark$sql$catalyst$expressions$codegen$CodeGenerator$$anon$1.load(CodeGenerator.scala:100)
  at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$$anon$1.load(CodeGenerator.scala:100)
  at org.sparkproject.guava.cache.LocalCache$LoadingValueReference.loadFuture(LocalCache.java:2539)
  at org.sparkproject.guava.cache.LocalCache$Segment.loadSync(LocalCache.java:2379)
  at org.sparkproject.guava.cache.LocalCache$Segment.lockedGetOrLoad(LocalCache.java:2342)
  at org.sparkproject.guava.cache.LocalCache$Segment.get(LocalCache.java:2257)
  at org.sparkproject.guava.cache.LocalCache.get(LocalCache.java:4000)
  at org.sparkproject.guava.cache.LocalCache.getOrLoad(LocalCache.java:4004)
  at org.sparkproject.guava.cache.LocalCache$LocalLoadingCache.get(LocalCache.java:4874)
  at org.apache.spark.sql.catalyst.expressions.codegen.CodeGenerator$.compile(CodeGenerator.scala:100)
  at org.apache.spark.sql.execution.WholeStageCodegenExec.liftedTree1$1(WholeStageCodegenExec.scala:73)
  at org.apache.spark.sql.execution.WholeStageCodegenExec.doExecute(WholeStageCodegenExec.scala:73)
  at org.apache.spark.sql.execution.SparkPlan.$anonfun$execute$1(SparkPlan.scala:194)
  at org.apache.spark.sql.execution.SparkPlan.$anonfun$executeQuery$1(SparkPlan.scala:232)
  at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:151)
  at org.apache.spark.sql.execution.SparkPlan.executeQuery(SparkPlan.scala:229)
  at org.apache.spark.sql.execution.SparkPlan.execute(SparkPlan.scala:190)
  at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.inputRDD$lzycompute(ShuffleExchangeExec.scala:100)
  at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.inputRDD(ShuffleExchangeExec.scala:100)
  at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.mapOutputStatisticsFuture$lzycompute(ShuffleExchangeExec.scala:100)
  at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.mapOutputStatisticsFuture(ShuffleExchangeExec.scala:100)

```

```

at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.$anonfun$submitShuffleJob$1
at org.apache.spark.sql.execution.SparkPlan.$anonfun$executeQuery$1(SparkPlan.scala:232)
at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:151)
at org.apache.spark.sql.execution.SparkPlan.executeQuery(SparkPlan.scala:229)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.submitShuffleJob(ShuffleExchangeLike.scala:100)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeLike.submitShuffleJob$(ShuffleExchangeLike.scala:100)
at org.apache.spark.sql.execution.exchange.ShuffleExchangeExec.submitShuffleJob(ShuffleExchangeExec.scala:100)
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.shuffleFuture$lzycompute(QueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.shuffleFuture(QueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.ShuffleQueryStageExec.doMaterialize(QueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.QueryStageExec.materialize(QueryStageExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysicalPlan$1(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysicalPlan$2(AdaptiveSparkPlanExec.scala:100)
at scala.collection.Iterator.foreach(Iterator.scala:943)
at scala.collection.Iterator.foreach$(Iterator.scala:943)
at scala.collection.AbstractIterator.foreach(Iterator.scala:1431)
at scala.collection.IterableLike.foreach(IterableLike.scala:74)
at scala.collection.IterableLike.foreach$(IterableLike.scala:73)
at scala.collection.AbstractIterable.foreach(Iterable.scala:56)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.$anonfun$getFinalPhysicalPlan$3(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.SparkSession.withActive(SparkSession.scala:779)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.getFinalPhysicalPlan(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.withFinalPlanUpdate(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.execution.adaptive.AdaptiveSparkPlanExec.executeCollect(AdaptiveSparkPlanExec.scala:100)
at org.apache.spark.sql.Dataset.$anonfun$collectToPython$1(Dataset.scala:3688)
at org.apache.spark.sql.Dataset.$anonfun$withAction$2(Dataset.scala:3858)
at org.apache.spark.sql.execution.QueryExecution$.withInternalError(QueryExecution.scala:51)
at org.apache.spark.sql.Dataset.$anonfun$withAction$1(Dataset.scala:3856)
at org.apache.spark.sql.execution.SQLExecution$. $anonfun$withNewExecutionId$6(SQLExecution.scala:100)
at org.apache.spark.sql.execution.SQLExecution$.withSQLConfPropagated(SQLExecution.scala:100)
at org.apache.spark.sql.execution.SQLExecution$. $anonfun$withNewExecutionId$1(SQLExecution.scala:100)
at org.apache.spark.sql.SparkSession.withActive(SparkSession.scala:779)
at org.apache.spark.sql.execution.SQLExecution$.withNewExecutionId(SQLExecution.scala:64)
at org.apache.spark.sql.Dataset.withAction(Dataset.scala:3856)
at org.apache.spark.sql.Dataset.collectToPython(Dataset.scala:3685)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:566)
at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
at py4j.Gateway.invoke(Gateway.java:282)
at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
at py4j.commands.CallCommand.execute(CallCommand.java:79)
at py4j.ClientServerConnection.waitForCommands(ClientServerConnection.java:182)
at py4j.ClientServerConnection.run(ClientServerConnection.java:106)
at java.base/java.lang.Thread.run(Thread.java:829)
Caused by: org.codehaus.janino.InternalCompilerException: Code of method "hashAgg_doAggregateWithPartitions"
at org.codehaus.janino.CodeContext.makeSpace(CodeContext.java:1051)
at org.codehaus.janino.CodeContext.write(CodeContext.java:947)

```

```

at org.codehaus.janino.CodeContext.writeShort(CodeContext.java:1072)
at org.codehaus.janino.UnitCompiler.writeConstantInterfaceMethodrefInfo(UnitCompiler.java:1
at org.codehaus.janino.UnitCompiler.invoke(UnitCompiler.java:11879)
at org.codehaus.janino.UnitCompiler.compileGet2(UnitCompiler.java:5186)
at org.codehaus.janino.UnitCompiler.access$9100(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$16.visitMethodInvocation(UnitCompiler.java:4482)
at org.codehaus.janino.UnitCompiler$16.visitMethodInvocation(UnitCompiler.java:4455)
at org.codehaus.janino.Java$MethodInvocation.accept(Java.java:5286)
at org.codehaus.janino.UnitCompiler.compileGet(UnitCompiler.java:4455)
at org.codehaus.janino.UnitCompiler.compileGetValue(UnitCompiler.java:5683)
at org.codehaus.janino.UnitCompiler.compileBoolean2(UnitCompiler.java:4175)
at org.codehaus.janino.UnitCompiler.access$6600(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$14.visitBinaryOperation(UnitCompiler.java:4008)
at org.codehaus.janino.UnitCompiler$14.visitBinaryOperation(UnitCompiler.java:3986)
at org.codehaus.janino.Java$BinaryOperation.accept(Java.java:5077)
at org.codehaus.janino.UnitCompiler.compileBoolean(UnitCompiler.java:3986)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:2474)
at org.codehaus.janino.UnitCompiler.access$1900(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitIfStatement(UnitCompiler.java:1498)
at org.codehaus.janino.UnitCompiler$6.visitIfStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$IfStatement.accept(Java.java:3140)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:2475)
at org.codehaus.janino.UnitCompiler.access$1900(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitIfStatement(UnitCompiler.java:1498)
at org.codehaus.janino.UnitCompiler$6.visitIfStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$IfStatement.accept(Java.java:3140)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:2486)
at org.codehaus.janino.UnitCompiler.access$1900(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$6.visitIfStatement(UnitCompiler.java:1498)
at org.codehaus.janino.UnitCompiler$6.visitIfStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java$IfStatement.accept(Java.java:3140)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)

```

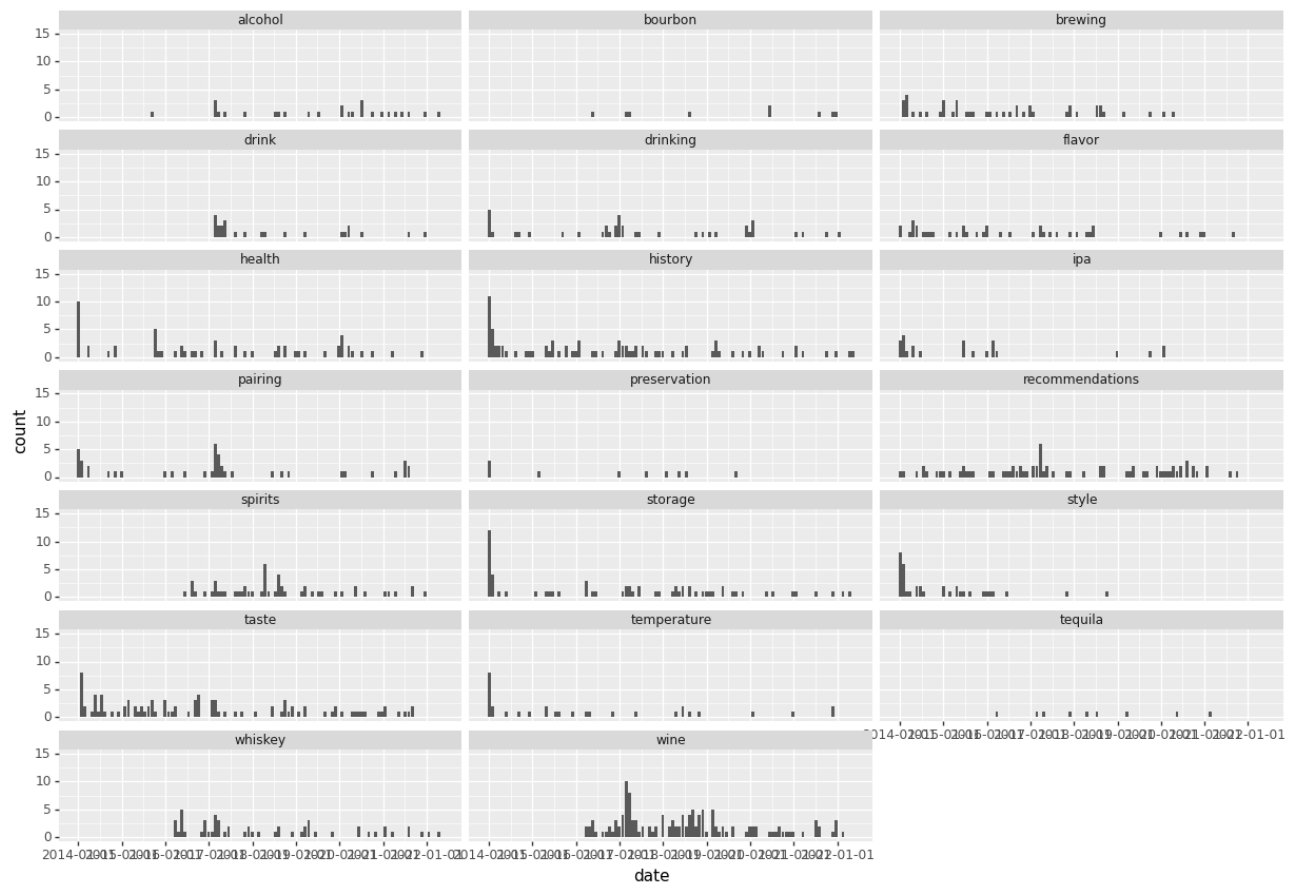
at org.codehaus.janino.UnitCompiler.access\$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1598)
at org.codehaus.janino.UnitCompiler.access\$2600(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitDoStatement(UnitCompiler.java:1505)
at org.codehaus.janino.UnitCompiler\$6.visitDoStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$DoStatement.accept(Java.java:3664)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access\$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1661)
at org.codehaus.janino.UnitCompiler.access\$2000(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitForStatement(UnitCompiler.java:1499)
at org.codehaus.janino.UnitCompiler\$6.visitForStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$ForStatement.accept(Java.java:3187)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1559)
at org.codehaus.janino.UnitCompiler.access\$1700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1496)
at org.codehaus.janino.UnitCompiler\$6.visitBlock(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$Block.accept(Java.java:2969)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:1848)
at org.codehaus.janino.UnitCompiler.access\$2200(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$6.visitWhileStatement(UnitCompiler.java:1501)
at org.codehaus.janino.UnitCompiler\$6.visitWhileStatement(UnitCompiler.java:1490)
at org.codehaus.janino.Java\$WhileStatement.accept(Java.java:3245)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:1490)
at org.codehaus.janino.UnitCompiler.compileStatements(UnitCompiler.java:1573)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:3420)
at org.codehaus.janino.UnitCompiler.compileDeclaredMethods(UnitCompiler.java:1362)
at org.codehaus.janino.UnitCompiler.compileDeclaredMethods(UnitCompiler.java:1335)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:807)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:975)
at org.codehaus.janino.UnitCompiler.access\$700(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler\$2.visitMemberClassDeclaration(UnitCompiler.java:392)
at org.codehaus.janino.UnitCompiler\$2.visitMemberClassDeclaration(UnitCompiler.java:384)
at org.codehaus.janino.Java\$MemberClassDeclaration.accept(Java.java:1445)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:384)
at org.codehaus.janino.UnitCompiler.compileDeclaredMemberTypes(UnitCompiler.java:1312)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:833)


```

at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:410)
at org.codehaus.janino.UnitCompiler.access$400(UnitCompiler.java:226)
at org.codehaus.janino.UnitCompiler$2.visitPackageMemberClassDeclaration(UnitCompiler.java:
at org.codehaus.janino.UnitCompiler$2.visitPackageMemberClassDeclaration(UnitCompiler.java:
at org.codehaus.janino.Java$PackageMemberClassDeclaration.accept(Java.java:1594)
at org.codehaus.janino.UnitCompiler.compile(UnitCompiler.java:384)
at org.codehaus.janino.UnitCompiler.compile2(UnitCompiler.java:362)
... 80 more
23/01/06 14:35:23 WARN WholeStageCodegenExec: Whole-stage codegen disabled for plan (id=4):
  *(4) HashAggregate(keys=[id#10, tag#34518, creation_date#8], functions=[], output=[id#10, tag#
+- *(4) Project [id#10, tag#34518, creation_date#8]
  +- *(4) BroadcastHashJoin [tag#34518], [tag#23691], Inner, BuildRight, false
    :- *(4) Filter isnotnull(tag#34518)
    :   +- *(4) Generate explode(tags_split#29166), [id#10, creation_date#8], false, [tag#3451
    :     +- *(4) Project [id#10, creation_date#8, CASE WHEN (size(tags_split#29062, true) =
    :       +- *(4) Project [id#10, creation_date#8, array_remove(array_distinct(array(regex
    :         +- *(4) Filter ((isnotnull(creation_date#8) AND (size(CASE WHEN (size(array_r
    :           +- *(4) ColumnarToRow
    :             +- FileScan parquet [creation_date#8,id#10,tags#19] Batched: true, Data
  +- BroadcastQueryStage 1
    +- BroadcastExchange HashedRelationBroadcastMode(List(input[0, string, false]),false)
    +- *(3) Filter isnotnull(tag#23691)
    +- TakeOrderedAndProject(limit=20, orderBy=[sum_views#23697L DESC NULLS LAST], c
    +- *(2) HashAggregate(keys=[tag#23691], functions=[sum(view_count#34545)], ou
    +- AQEShuffleRead coalesced
    +- ShuffleQueryStage 0
      +- Exchange hashpartitioning(tag#23691, 200), ENSURE_REQUIREMENTS, M
      +- *(1) HashAggregate(keys=[tag#23691], functions=[partial_sum(vi
      +- *(1) Project [tag#23691, view_count#34545]
      +- *(1) Generate explode(tags_split#18439), [view_count#345
      +- *(1) Project [view_count#34545, CASE WHEN (size(tags_
      +- *(1) Project [view_count#34545, array_remove(array
      +- *(1) Filter ((isnotnull(view_count#34545) AND (
      +- *(1) ColumnarToRow
      +- FileScan parquet [tags#34543,view_count#3

```

/config/workspace/.venv/lib/python3.10/site-packages/plotnine/layer.py:391: PlotnineWarning: po



```
<ggplot: (8788912962425)>
```

3.7 najczęściej pojawiające się słowa w tytułach (z pominięciem stopwords)

```
from bs4 import BeautifulSoup
from html import unescape
from pyspark.sql.functions import udf, regexp_replace
from pyspark.sql.types import *
# remove html tags
def tags_remove(s):
    soup = BeautifulSoup(unescape(s), 'lxml')
    return soup.text

udf_tags_remove = udf(lambda m: tags_remove(m))

titles = posts.filter(f.col("title").isNotNull()).select(f.col("title"))\
    .withColumn("title_clean", f.lower(f.col("title")))\
    .withColumn("title_clean", regexp_replace('title_clean', "[^a-zA-Z\\s]", " "))

from pyspark.ml.feature import Tokenizer, StopWordsRemover
```

```

from nltk.stem.snowball import SnowballStemmer
udf_filter_length = udf(lambda row: [x for x in row if len(x) > 1], ArrayType(StringType()))

stemmer = SnowballStemmer(language='english')
stemmer_udf = udf(lambda token: stemmer.stem(token), StringType())

tokenizer = Tokenizer(inputCol='title_clean', outputCol='words_token')
title_tokens = tokenizer.transform(titles).withColumn('words_token', udf_filter_length(f.col('words_token')))

remover = StopWordsRemover(inputCol='words_token', outputCol='words_no_stop')
title_tokens_no_stop = remover.transform(title_tokens)
exploded = title_tokens_no_stop.withColumn("words", f.explode(f.col("words_no_stop")))

title_stem = exploded.withColumn('words_stem', stemmer_udf("words"))

word_lookup = title_stem.select([f.col("words"), f.col("words_stem")]).distinct()
word_lookup.show() # TODO aggregate this
title_stem.groupBy("words_stem").agg(f.count("title").alias('count')).orderBy('count', ascen

```

```

+-----+-----+
|      words| words_stem|
+-----+-----+
|      opened|      open|
|antidepressants|antidepress|
|      taken|      taken|
|  alternative|  altern|
|    learning|    learn|
|    sherry|    sherri|
|  regionali| regionali|
|  archetype|  archetyp|
|  inhibitor|  inhibitor|
|    outside|    outsid|
|      bay|      bay|
|    sangria|    sangria|
|    invest|    invest|
|    together|    together|
|    fake|    fake|
|    kahlua|    kahlua|
|    imported|    import|
|    tables|    tabl|
|    desire|    desir|
|    bavaria|    bavaria|
+-----+-----+

```

only showing top 20 rows

```
+-----+-----+
|words_stem|count|
+-----+-----+
|      beer|  476|
|      wine|  147|
|     drink|  104|
|  alcohol|   88|
|    differ|   72|
|    bottl|   68|
|       use|   50|
|     tast|   47|
|     brew|   43|
|     make|   41|
|     good|   33|
| cocktail|   29|
|      age|   27|
| recommend|   26|
|      ale|   26|
|     like|   24|
|     made|   23|
|  whiskey|   23|
|   spirit|   23|
|      one|   22|
+-----+-----+
```

only showing top 20 rows

3.8 procent użytkowników którzy nigdy nic nie zapostowali

```
# users.show(2)
# posts.show(2)

users_posts = (users.join(
  (posts.select(f.col('id').alias('post_id'), f.col('owner_user_id'))), on=[users.id == po
  .filter("id IS NOT NULL and NOT id = -1")
))

users_posts.select([f.col('id'), f.col('post_id')])\
  .groupBy("id")\
  .agg(
    f.count(f.col('post_id')).alias('post_count')
  )\
  .agg(
    f.sum(f.when(f.col('post_count') == 0, f.lit(1)).otherwise(f.lit(0))).alias("not_pos
    f.count('id').alias('all')
  ) \
```

```

        .withColumn('% not posted', (f.col("not_posted") / f.col('all') * 100)).show()

# (users.join(
#     (posts.select(f.col('id').alias('post_id'), f.col('owner_user_id'))), on=[users.id ==
#     .filter(f.col('id').isNotNull() & f.col('post_id').isNull())\
#     .select([f.col('id'), f.col('display_name')]).distinct()\
#     .show()
# )

```

```

+-----+-----+-----+
|not_posted| all|      % not posted|
+-----+-----+-----+
|      7691|8947|85.96177489661339|
+-----+-----+-----+

```

3.9 średni czas od pojawienia się pytania do pojawienia się zaakceptowanej odpowiedzi

```

# keep only questions with answers
questions = posts.filter(f.col('post_type_id') == 1).filter(f.col('answer_count') > 0).select(
answers = posts.filter(f.col('post_type_id') == 2).select([f.col('id').alias('a_id'), f.col('answer_count')])
#posts.show(1, vertical=True)
time_to_accept = questions.join(answers, on=[questions.accepted_answer_id==answers.a_id])\
    .withColumn('time_to_accept_sec', f.unix_timestamp('a_creation_date') - f.unix_timestamp('q_creation_date'))\
    .withColumn('time_to_accept_min', f.round(f.col('time_to_accept_sec') / 60, 2))

time_to_accept.agg(
    f.avg('time_to_accept_min'),
    f.stddev('time_to_accept_min'),
    f.percentile_approx("time_to_accept_min", [0.25, 0.5, 0.75], 1000000).alias("quantiles")
).show(truncate=False)

time_to_accept_pd = time_to_accept.withColumn('all', f.lit("all_questions")).toPandas()

from plotnine import geom_jitter
(ggplot(time_to_accept_pd, aes(x='all', y="time_to_accept_min"))\
+geom_jitter())


```

```

+-----+-----+-----+
|avg(time_to_accept_min)|stddev_samp(time_to_accept_min)|quantiles|
+-----+-----+-----+
|25244.9435714286      |123338.6325101642      |[141.53, 753.25, 3605.8]|
+-----+-----+-----+

```

/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: FutureWarning: The default value of chunksize will be 1000000 in a future version of pandas. Specify a value to silence this warning.



<ggplot: (8788909444688)>

3.9.1 remove outliers


```
no_outliers = time_to_accept.filter(f.col('time_to_accept_min') < 100000)

no_outliers.agg(
    f.avg('time_to_accept_min'),
    f.stddev('time_to_accept_min'),
    f.percentile_approx("time_to_accept_min", [0.25, 0.5, 0.75], 1000000).alias("quantil
).show(truncate=False)

no_outliers_pd = no_outliers.withColumn('all', f.lit("all_questions")).toPandas()
(ggplot(no_outliers_pd, aes(x='all', y="time_to_accept_min"))\
+geom_jitter())
```

avg(time_to_accept_min)	stddev_samp(time_to_accept_min)	quantiles
4769.308792048925	12694.627623679486	[127.72, 644.93, 2922.82]

```
/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: Futu
/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: Futu
```



```
<ggplot: (8788909380107)>
```

References