

Wieloskalowa analiza danych z forum internetowego przy użyciu usług chmury AWS

Michał Kamiński

21.01.2023

Spis treści

Streszczenie	4
Summary	4
 I Wstęp	 5
Technologie big data	6
Formaty danych	6
Chmury	6
Cel pracy	6
 II Wyniki i Dyskusja	 7
 1 Schemat infrastruktury	 8
1.1 Ekstrakcja	8
1.2 Przygotowanie danych wstępnych	9
1.3 Budowa infrastruktury	9
 2 Wstępna obróbka danych	 10
2.1 Konfiguracja aplikacji	10
2.2 Schematy danych	11
2.2.1 Users	11
2.2.2 Tags	12
2.2.3 Votes	12
2.2.4 Posts	13
2.2.5 Post links	14
2.2.6 Post History	15
2.2.7 Badges	16
2.3 Czyszczenie kolumn tekstowych	16
2.4 Zapis danych jako plik w formacie parquet	19
 3 Eksploracja i analiza danych	 22
3.1 Questions/Answers over time	22
3.2 Forum retention - time from account creation to last comments	24
3.3 porównanie najwyżej i najniżej ocenianych pytań (długość, tagi, liczba odpowiedzi)	27
3.4 procent przypadków kiedy najwyżej oceniana odpowiedź to nie zaakceptowana odpowiedź	36
3.5 rozkład ocen odpowiedzi zaakceptowanych vs pozostałych (średnia, odchylenie, minimum, maksimum)	37
3.6 top N tagów które wygenerowały najwięcej wyświetleń	38
3.7 liczba postów w czasie dla każdego z top N tagów (lineplot/barplot)	39
3.8 najczęściej pojawiające się słowa w tytułach (z pominięciem stopwords)	39
3.9 procent użytkowników którzy nigdy nic nie zapostowali	41

3.10	średni czas od pojawienia się pytania do pojawienia się zaakceptowanej odpowiedzi . . .	42
3.10.1	remove outliers	43
References		45
4	Załączniki	46
4.1	Polecenia budujące infrastrukturę	46
4.1.1	EMR	46
4.1.2	S3	46

Streszczenie

Słowa kluczowe: Big Data, Spark, AWS, EMR, S3

Summary

Keywords: Big Data, Spark, AWS, EMR, S3

Cześć I

Wstęp

Technologie big data

Formaty danych

Chmury

Cel pracy

Celem niniejszej pracy jest utworzenie infrastruktury w chmurze obliczeniowej AWS pozwalające na wielkoskalową analizy danych w sytemie rozproszonym (ang. *Big Data*).

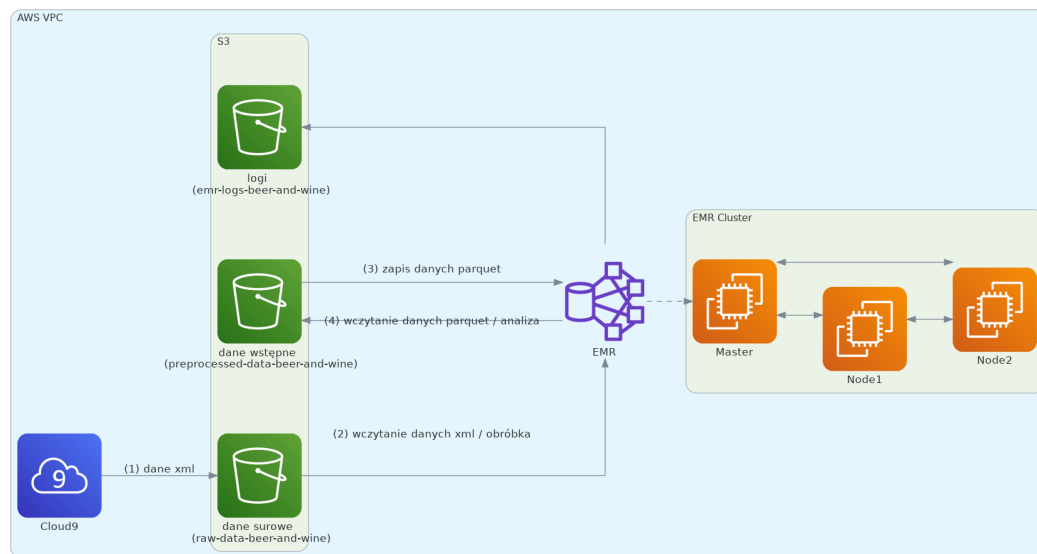
Do stworzenia przykładowego projektu wykorzystano dane ze strony [Stack Exchange](#) zawierającej zestawy danych pochodzące z forów społecznościowych. Analizę ograniczono do danych pochodzących z forum o nazwie [Beer, Wine and Spirits](#).

W niniejszej pracy ...

Cześć II

Wyniki i Dyskusja

1 Schemat infrastruktury



Rysunek 1.1: Schemat rozwiązania

W celu rozwiązania postawionego problemu analitycznego stworzono infrastrukturę wyłącznie w obrębie chmury AWS, której ogólny schemat przedstawiono na Rysunek 1.1

1.1 Ekstrakcja

Do etapu ekstrakcji danych wykorzystano usługę Cloud9, która zapewnia dostęp do terminala maszyny wirtualnej z systemem linux (platforma Amazon Linux 2, typ instancji t2.micro). Z użyciem tej usługi dane zostały pobrane ze źródła w binarnym formacie 7z a następnie pliki zostały wyekstrahowane w formacie xml przy pomocy programu p7zip. Dane w formacie xml zostały następnie skopiowane do serwisu S3, gdzie utworzono koszyk danych (ang. *bucket*) o nazwie **raw-data-beer-and-wine**, którego przeznaczeniem jest przetwarzanie danych nieprzetworzonych.

Powyższe operacje zostały wykonane przy użyciu poniższych poleceń:


```
# instalacja programu p7zip
sudo yum install p7zip.x86_64

# pobranie danych
wget https://archive.org/download/stackexchange/beer.stackexchange.com.7z

# ekstrakcja danych do folderu raw-data
7za e beer.stackexchange.com.7z -oraw-data

# zapis danych do koszyka S3 przy użyciu programu `AWS CLI`
aws s3 cp $(pwd)/raw-data s3://raw-data-beer-and-wine/ --recursive --include "*.xml"
```

1.2 Przygotowanie danych wstępnych

W celu przygotowania danych do analizy, dane surowe zostały wstępnie przetworzone oraz zapisane w formacie **parquet**, co pozwoli na wydajniejsze wczytywanie danych podczas uruchomień programu. Podczas etapu wstępnego przetwarzania danych, oprócz zmiany formatu plików, zdefiniowane zostały także schematy danych, które zapewniają, że kolumny danych będą posiadały odpowiednie typy oraz, że krytyczne dane nie będą zawierały pustych wartości. Dodatkowo kolumny z wartościami tekstowymi, niesłownikowanymi zostały oczyszczone z tagów **html** oraz poddane standardowej procedurze oczyszczania tekstu.

Powyższe czynności zostały wykonane w notatniku typu Jupiter (ang. *Jupyter Notebook*) w serwisie **AWS EMR**. Stworzono klaster EMR (wersja 6.8.0) z instalacją Hadoop 3.2.1, Jupyter Hub oraz Spark 3.3.0, składający się z 1 instancji typu *master* oraz 2 instancji typu *core*, każda typu **m4.large**. W celu ograniczenia kosztów jako opcję zakupu wybrano typ **spot** z limitem maksymalnym ceny odpowiadającej typowi **on-demand**. Wielkość dysków EBS stworzonych instancji wynosiła 32 GiB dla każdej instancji w klastrze.

Polecenie programu AWS CLI odpowiadające za utworzenie klastra znajduje się w sekcji Sekcja 4.1.1.

Dostęp do **Jupyter Notebook** w utworzonym klastrze jest możliwy poprzez połączenie przez przeglądarkę z środowiskiem graficznym Jupyter Hub wykorzystując adres DNS instancji *master* i port 9443.

1.3 Budowa infrastruktury

Wszystkie serwisy AWS na potrzeby tego projektu zostały utworzone w sposób programatyczny przy użyciu programu AWS CLI (poza Cloud9, który został utworzony z poziomu konsoli zarządzającej). Wykorzystane polecenia dostępne są w sekcji Rozdział 4.

2 Wstępna obróbka danych

2.1 Konfiguracja aplikacji

W celu przygotowania danych do analizy zostały one wstępnie przetworzone. Pierwszym etapem wstępnego przetwarzania jest wczytanie danych do środowiska analitycznego. Dane surowe, przechowywane w koszyku `raw-data-beer-and-wine` znajdowały się w mało przyjaznym dla analiz formacie `xml`. Wczytanie tego typu danych wymagało załadowania dodatkowego pakietu `jar` o nazwie `spark-xml_2.12:0.14.0` pobranego z repozytorium `maven`.

W serwisie EMR można dodać tego typu pakiety wykorzystując specjalne polecenia typu `Sparkmagic` rozpoczynające się od znaków `%%`. W tym przypadku użyto `%%configure`:

```
%%configure -f
{
  "conf": {
    "spark.jars.packages": "com.databricks:spark-xml_2.12:0.14.0"
  }
}

from pyspark.sql import SparkSession
spark = SparkSession.builder.appName("Preprocessing").getOrCreate()
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	User	Current session?
2	application_1674145937048_0003	pyspark	idle	Link	Link	None	

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

SparkSession available as 'spark'.

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

2.2 Schematy danych

2.2.1 Users

```
from pyspark.sql.types import *

users_schema = StructType([
    StructField('_AboutMe', StringType(), True),
    StructField('_AccountId', IntegerType(), True),
    StructField('_CreationDate', TimestampType(), True),
    StructField("_DisplayName", StringType(), True),
    StructField("_DownVotes", IntegerType(), True),
    StructField("_Id", IntegerType(), True),
    StructField("_LastAccessDate", TimestampType()),
    StructField("_Location", StringType(), True),
    StructField("_ProfileImageUrl", StringType(), True),
    StructField("_Reputation", IntegerType(), True),
    StructField("_UpVotes", IntegerType(), True),
    StructField("_Views", IntegerType(), True),
    StructField("_WebsiteUrl", StringType(), True)
])

users = spark.read.format('xml').options(rowTag='row').schema(users_schema).load("s3://raw-d
users.show(2, vertical=True, truncate=50)
```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px

```
-RECORD 0-----
 _AboutMe      | <p>Hi, I'm not really a person.</p>\n\n<p>I'm a...
 _AccountId    | -1
 _CreationDate | 2014-01-21 17:45:53.587
 _DisplayName  | Community
 _DownVotes    | 503
 _Id           | -1
 _LastAccessDate | 2014-01-21 17:45:53.587
 _Location     | on the server farm
 _ProfileImageUrl | null
 _Reputation   | 1
 _UpVotes      | 2
 _Views        | 5
 _WebsiteUrl   | http://meta.stackexchange.com/
-RECORD 1-----
 _AboutMe      | <p>Dev #2 who helped create Stack Overflow curr...
 _AccountId    | 2
 _CreationDate | 2014-01-21 20:21:18.797
 _DisplayName  | Geoff Dalgas
 _DownVotes    | 0
 _Id           | 1
```

```

_LastAccessDate | 2016-05-06 20:34:57.983
_Location       | Corvallis, OR
_ProfileImageUrl | https://i.stack.imgur.com/nD1lk.png?s=128&g=1
_Reputation     | 101
_UpVotes        | 0
_Views          | 42
_WebsiteUrl     | http://stackoverflow.com

```

only showing top 2 rows

2.2.2 Tags

```

tags_schema = StructType([
    StructField('_Count', IntegerType(), True),
    StructField('_ExcerptPostId', IntegerType(), True),
    StructField('_Id', IntegerType(), True),
    StructField("_TagName", StringType(), True),
    StructField("_WikiPostId", IntegerType(), True)
])

tags = spark.read.format('xml').options(rowTag='row').schema(tags_schema).load("s3://raw-dat
tags.show(n=5)

```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px

```

+-----+-----+-----+-----+-----+
|_Count|_ExcerptPostId|_Id|_TagName|_WikiPostId|
+-----+-----+-----+-----+-----+
|    17|          5062| 1|    hops|          5061|
|    85|          7872| 2|   history|          7871|
|    69|          4880| 4|   brewing|          4879|
|    37|          5109| 5|   serving|          5108|
|    31|           304| 6|temperature|           303|
+-----+-----+-----+-----+-----+

```

only showing top 5 rows

2.2.3 Votes

```

votes_schema = StructType([
    StructField('_BountyAmount', IntegerType(), True),
    StructField('_CreationDate', TimestampType(), True),
    StructField('_Id', IntegerType(), True),
    StructField("_PostId", StringType(), True),
    StructField("_UserId", IntegerType(), True),
    StructField("_VoteTypeId", IntegerType(), True)
])

votes = spark.read.format('xml').options(rowTag='row').schema(votes_schema).load("s3://raw-d

```

```
votes.show(n=5)
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

```
+-----+-----+-----+-----+-----+
|_BountyAmount|_CreationDate|_Id|_PostId|_UserId|_VoteTypeId|
+-----+-----+-----+-----+-----+
|          null|2014-01-21 00:00:00| 1|      1|   null|          2|
|          null|2014-01-21 00:00:00| 2|      1|   null|          2|
|          null|2014-01-21 00:00:00| 3|      4|   null|          2|
|          null|2014-01-21 00:00:00| 4|      1|   null|          2|
|          null|2014-01-21 00:00:00| 5|      4|   null|          2|
+-----+-----+-----+-----+-----+
```

only showing top 5 rows

2.2.4 Posts

```
posts_schema = StructType([
    StructField('_AcceptedAnswerId', IntegerType(), True),
    StructField('_AnswerCount', IntegerType(), True),
    StructField('_Body', StringType(), True),
    StructField("_ClosedDate", TimestampType(), True),
    StructField("_CommentCount", IntegerType(), True),
    StructField("_CommunityOwnedDate", TimestampType(), True),
    StructField("_ContentLicense", StringType(), True),
    StructField("_CreationDate", TimestampType(), True),
    StructField("_FavoriteCount", IntegerType(), True),
    StructField("_Id", IntegerType(), True),
    StructField("_LastActivityDate", TimestampType(), True),
    StructField("_LastEditDate", TimestampType(), True),
    StructField("_LastEditorDisplayName", StringType(), True),
    StructField("_LastEditorUserId", IntegerType(), True),
    StructField("_OwnerDisplayName", StringType(), True),
    StructField("_OwnerUserId", IntegerType(), True),
    StructField("_ParentId", IntegerType(), True),
    StructField("_PostTypeId", IntegerType(), True),
    StructField("_Score", IntegerType(), True),
    StructField("_Tags", StringType(), True),
    StructField("_Title", StringType(), True),
    StructField("_ViewCount", IntegerType(), True),
])
```

```
posts = spark.read.format('xml').options(rowTag='row').schema(posts_schema).load("s3://raw-d
posts.show(n=1,vertical=True, truncate=50)
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

```

-RECORD 0-----
|_AcceptedAnswerId      | 4
|_AnswerCount           | 1
|_Body                  | <p>I was offered a beer the other day that was ...
|_ClosedDate            | null
|_CommentCount          | 0
|_CommunityOwnedDate    | null
|_ContentLicense        | CC BY-SA 3.0
|_CreationDate          | 2014-01-21 20:26:05.383
|_FavoriteCount         | null
|_Id                    | 1
|_LastActivityDate       | 2014-01-21 22:04:34.977
|_LastEditDate          | 2014-01-21 22:04:34.977
|_LastEditorDisplayName | null
|_LastEditorUserId      | 8
|_OwnerDisplayName       | null
|_OwnerUserId           | 7
|_ParentId              | null
|_PostTypeId            | 1
|_Score                 | 21
|_Tags                  | <hops>
|_Title                 | What is a citra hop, and how does it differ fro...
|_ViewCount             | 2441
only showing top 1 row

```

2.2.5 Post links

```

links_schema = StructType([
    StructField("_CreationDate", TimestampType()),
    StructField("_Id", IntegerType()),
    StructField("_LinkTypeId", IntegerType()),
    StructField("_PostId", IntegerType()),
    StructField("_RelatedPostId", IntegerType())
])

links = spark.read.format('xml').options(rowTag='row').schema(links_schema).load("s3://raw-d
links.show(n=5, truncate=False)

```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px'
```

```

+-----+-----+-----+-----+-----+
|_CreationDate      |_Id|_LinkTypeId|_PostId|_RelatedPostId|
+-----+-----+-----+-----+-----+
|2014-01-21 21:04:25.23|25|3          |29      |25          |
|2014-01-21 21:42:09.103|89|1          |83      |50          |
|2014-01-21 21:50:41.313|95|1          |86      |2           |
|2014-01-21 22:07:35.783|101|3         |47      |99          |
|2014-01-21 22:13:51.38 |102|1         |74      |3           |

```

```
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

2.2.6 Post History

```
history_schema = StructType([
    StructField("_Comment", StringType()),
    StructField("_ContentLicense", StringType()),
    StructField("_CreationDate", TimestampType()),
    StructField("_Id", IntegerType()),
    StructField("_PostHistoryTypeId", IntegerType()),
    StructField("_PostId", IntegerType()),
    StructField("_RevisionGUID", StringType()),
    StructField("_Text", StringType()),
    StructField("_UserDisplayName", StringType()),
    StructField("_UserId", IntegerType()),
])

history = spark.read.format('xml').options(rowTag='row').schema(history_schema).load("s3://r
history.show(n=2,vertical=True, truncate=50)
```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px

```
-RECORD 0-----
 _Comment      | null
 _ContentLicense | CC BY-SA 3.0
 _CreationDate  | 2014-01-21 20:26:05.383
 _Id           | 1
 _PostHistoryTypeId | 2
 _PostId       | 1
 _RevisionGUID  | a17002a0-00b0-417b-a404-0d8864bbbca5
 _Text         | I was offered a beer the other day that was rep...
 _UserDisplayName | null
 _UserId       | 7
-RECORD 1-----
 _Comment      | null
 _ContentLicense | CC BY-SA 3.0
 _CreationDate  | 2014-01-21 20:26:05.383
 _Id           | 2
 _PostHistoryTypeId | 1
 _PostId       | 1
 _RevisionGUID  | a17002a0-00b0-417b-a404-0d8864bbbca5
 _Text         | What is a citra hop, and how does it differ fro...
 _UserDisplayName | null
 _UserId       | 7
only showing top 2 rows
```

2.2.7 Badges

```
badges_schema = StructType([
    StructField("_Class", IntegerType()),
    StructField("_Date", TimestampType()),
    StructField("_Id", IntegerType()),
    StructField("_Name", StringType()),
    StructField("_TagBased", BooleanType()),
    StructField("_UserId", IntegerType()),
])

badges = spark.read.format('xml').options(rowTag='row').schema(badges_schema).load("s3://raw
badges.show(n=5,truncate=False)
```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px

_Class	_Date	_Id	_Name	_TagBased	_UserId
3	2014-01-21 20:52:16.97	1	Autobiographer	false	1
3	2014-01-21 20:52:16.97	2	Autobiographer	false	2
3	2014-01-21 20:52:16.97	3	Autobiographer	false	6
3	2014-01-21 20:52:16.97	4	Autobiographer	false	7
3	2014-01-21 20:52:16.97	5	Autobiographer	false	9

only showing top 5 rows

2.3 Czyszczenie kolumn tekstowych

```
from pyspark.sql.functions import regexp_replace, trim, udf, col

from bs4 import BeautifulSoup
from html import unescape

def tags_remove(s):
    if s is not None:
        soup = BeautifulSoup(unescape(s), 'lxml')
        return soup.text
    else:
        return None

udf_tags_remove = udf(lambda m: tags_remove(m))
```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px


```

users_clean = users.withColumn("_AboutMe_clean", regexp_replace("_AboutMe", "\n|\t|\r", " "))
                  .withColumn("_AboutMe_clean", udf_tags_remove(col('_AboutMe_clean'))) \
                  .withColumn("_AboutMe_clean", regexp_replace("_AboutMe_clean", "\s{2,}", " ")) \
                  .withColumn("_AboutMe_clean", trim("_AboutMe_clean"))

history_clean = history.withColumn("_Text_clean", regexp_replace("_Text", "\n|\t|\r", " "))
                      .withColumn("_Text_clean", udf_tags_remove(col('_Text_clean'))) \
                      .withColumn("_Text_clean", regexp_replace("_Text_clean", "\s{2,}", " ")) \
                      .withColumn("_Text_clean", trim("_Text_clean"))

posts_clean = posts.withColumn("_Body_clean", regexp_replace("_Body", "\n|\t|\r", " ")) \
                    .withColumn("_Body_clean", udf_tags_remove(col('_Body_clean'))) \
                    .withColumn("_Body_clean", regexp_replace("_Body_clean", "\s{2,}", " ")) \
                    .withColumn("_Body_clean", trim("_Body_clean"))

```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

```
users_clean.show(1, vertical=True, truncate=50)
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

```

-RECORD 0-----
 _AboutMe      | <p>Hi, I'm not really a person.</p>\n\n<p>I'm a...
 _AccountId    | -1
 _CreationDate | 2014-01-21 17:45:53.587
 _DisplayName   | Community
 _DownVotes    | 503
 _Id           | -1
 _LastAccessDate | 2014-01-21 17:45:53.587
 _Location     | on the server farm
 _ProfileImageUrl | null
 _Reputation   | 1
 _UpVotes      | 2
 _Views        | 5
 _WebsiteUrl   | http://meta.stackexchange.com/
 _AboutMe_clean | Hi, I'm not really a person. I'm a background p...
only showing top 1 row

```

```
history_clean.show(1, vertical=True, truncate=50)
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

```

-RECORD 0-----
 _Comment      | null

```

```

_ContentLicense      | CC BY-SA 3.0
_CreationDate        | 2014-01-21 20:26:05.383
_Id                  | 1
_PostHistoryTypeId   | 2
_PostId              | 1
_RevisionGUID        | a17002a0-00b0-417b-a404-0d8864bbbca5
_Text                | I was offered a beer the other day that was rep...
_UserDisplayName      | null
_UserId              | 7
_Text_clean          | I was offered a beer the other day that was rep...
only showing top 1 row

```

```
posts_clean.show(1, vertical=True, truncate=50)
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

```

-RECORD 0-----
_AcceptedAnswerId    | 4
_AnswerCount         | 1
_Body                | <p>I was offered a beer the other day that was ...
_ClosedDate          | null
_CommentCount        | 0
_CommunityOwnedDate  | null
_ContentLicense      | CC BY-SA 3.0
_CreationDate        | 2014-01-21 20:26:05.383
_FavoriteCount       | null
_Id                  | 1
_LastActivityDate    | 2014-01-21 22:04:34.977
_LastEditDate        | 2014-01-21 22:04:34.977
_LastEditorDisplayName | null
_LastEditorUserId    | 8
_OwnerDisplayName     | null
_OwnerUserId         | 7
_ParentId            | null
_PostTypeId          | 1
_Score               | 21
_Tags                | <hops>
_Title               | What is a citra hop, and how does it differ fro...
_ViewCount           | 2441
_Body_clean          | I was offered a beer the other day that was rep...
only showing top 1 row

```

2.4 Zapis danych jako plik w formacie parquet

```
users_clean.select(
    col("_AboutMe").alias("about_me"),
    col("_AboutMe_clean").alias("about_me_clean"),
    col("_CreationDate").alias("creation_date"),
    col("_DisplayName").alias("display_name"),
    col("_DownVotes").alias("down_votes"),
    col("_Id").alias("id"),
    col("_LastAccessDate").alias("last_access_date"),
    col("_Location").alias("location"),
    col("_ProfileImageUrl").alias("profile_image_url"),
    col("_Reputation").alias("reputatio"),
    col("_UpVotes").alias("up_votes"),
    col("_Views").alias("views"),
    col("_WebsiteUrl").alias("website_url")
).write.mode('overwrite').format('parquet').option('path', "s3://preprocessed-data-beer-and-
```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px'

```
tags.select(
    col("_Count").alias("count"),
    col("_ExcerptPostId").alias("excerpt_post_id"),
    col("_Id").alias("id"),
    col("_TagName").alias("tag_name"),
    col("_WikiPostId").alias("wiki_post_id"),
).write.mode('overwrite').format('parquet').option('path', "s3://preprocessed-data-beer-and-
```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px'

```
votes.select(
    col("_BountyAmount").alias("bounty_amount"),
    col("_CreationDate").alias("creation_date"),
    col("_Id").alias("id"),
    col("_PostId").alias("post_id"),
    col("_UserId").alias("user_id"),
    col("_VoteTypeId").alias("vote_type_id"),
).write.mode('overwrite').format('parquet').option('path', "s3://preprocessed-data-beer-and-
```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px'

```
posts_clean.select(
    col("_AcceptedAnswerId").alias("accepted_answer_id"),
    col("_AnswerCount").alias("answer_count"),
```

```

col("_Body").alias("body"),
col("_Body_clean").alias("body_clean"),
col("_ClosedDate").alias("closed_date"),
col("_CommentCount").alias("comment_count"),
col("_CommunityOwnedDate").alias("community_owned_date"),
col("_ContentLicense").alias("content_licence"),
col("_CreationDate").alias("creation_date"),
col("_FavoriteCount").alias("favourite_count"),
col("_Id").alias("id"),
col("_LastActivityDate").alias("last_activity_date"),
col("_LastEditDate").alias("last_edit_date"),
col("_LastEditorDisplayName").alias("last_editor_display_name"),
col("_LastEditorUserId").alias("last_editor_user_id"),
col("_OwnerUserId").alias("owner_user_id"),
col("_PostTypeId").alias("post_type_id"),
col("_ParentId").alias("parent_id"),
col("_Score").alias("score"),
col("_Tags").alias("tags"),
col("_Title").alias("title"),
col("_ViewCount").alias("view_count"),
).write.mode('overwrite').format('parquet').option('path', "s3://preprocessed-data-beer-and-

```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px

```

links.select(
    col("_CreationDate").alias("creation_date"),
    col("_Id").alias("id"),
    col("_LinkTypeId").alias("link_type_id"),
    col("_PostId").alias("post_id"),
    col("_RelatedPostId").alias("related_post_id"),
).write.mode('overwrite').format('parquet').option('path', "s3://preprocessed-data-beer-and-

```

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px

```

history_clean.select(
    col("_Comment").alias("comment"),
    col("_ContentLicense").alias("content_license"),
    col("_CreationDate").alias("creation_date"),
    col("_Id").alias("id"),
    col("_PostHistoryTypeId").alias("post_history_type_id"),
    col("_PostId").alias("post_id"),
    col("_RevisionGUID").alias("revision_guid"),
    col("_Text").alias("text"),
    col("_Text_clean").alias("text_clean"),
    col("_UserDisplayName").alias("user_display_name"),
    col("_UserId").alias("user_id"),

```

```
).write.mode('overwrite').format('parquet').option('path', "s3://preprocessed-data-beer-and-
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

```
badges.select(  
    col("_Class").alias("class"),  
    col("_Date").alias("date"),  
    col("_Id").alias("id"),  
    col("_Name").alias("name"),  
    col("_TagBased").alias("tag_based"),  
    col("_UserId").alias("user_id"),  
).write.mode('overwrite').format('parquet').option('path', "s3://preprocessed-data-beer-and-
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px
```

3 Eksploracja i analiza danych

3.1 Questions/Answers over time

```
from pyspark.sql import (  
    SparkSession,  
    functions as f  
)  
import matplotlib
```

```
spark = SparkSession.builder.master("local[12]").appName("Analytics").getOrCreate()
```

```
posts = spark.read.format('parquet').load("outputs/posts")  
posts.show(1, vertical=True)
```

```
-RECORD 0-----  
accepted_answer_id      | 4  
answer_count            | 1  
body                    | <p>I was offered ...  
body_clean              | I was offered a b...  
closed_date             | null  
comment_count           | 0  
community_owned_date    | null  
content_licence         | CC BY-SA 3.0  
creation_date           | 2014-01-21 20:26:...  
favourite_count         | null  
id                      | 1  
last_activity_date      | 2014-01-21 22:04:...  
last_edit_date          | 2014-01-21 22:04:...  
last_editor_display_name | null  
last_editor_user_id     | 8  
owner_user_id           | 7  
post_type_id            | 1  
parent_id               | null  
score                   | 21  
tags                    | <hops>  
title                   | What is a citra h...  
view_count              | 2434  
only showing top 1 row
```

```
posts_grouped = (  
    posts  
    .filter(f.col('owner_user_id').isNotNull())
```

```

.groupBy(
  f.window('creation_date', '4 weeks')
)
.agg(
  f.sum(f.lit(1)).alias('all'),
  f.sum(f.when(f.col('post_type_id') == 1, f.lit(1)).otherwise(f.lit(0))).alias('questions'),
  f.sum(f.when(f.col('post_type_id') == 2, f.lit(1)).otherwise(f.lit(0))).alias('answers')
)
# window struct has nested columns 'start' and 'end'
.withColumn('date', f.col('window.start').cast('date'))
.orderBy('date')
).toPandas()

```

```
posts_grouped.head()
```

	window	all	questions	answers	date
0	(2014-01-02 00:00:00, 2014-01-30 00:00:00)	413	150	243	2014-01-02
1	(2014-01-30 00:00:00, 2014-02-27 00:00:00)	190	58	118	2014-01-30
2	(2014-02-27 00:00:00, 2014-03-27 00:00:00)	50	16	34	2014-02-27
3	(2014-03-27 00:00:00, 2014-04-24 00:00:00)	47	16	31	2014-03-27
4	(2014-04-24 00:00:00, 2014-05-22 00:00:00)	44	10	34	2014-04-24

```

# posts_grouped.plot(
#   x='date',
#   figsize=(12, 6),
#   title='Number of questions/answers per month (4 weeks)',
#   legend=True,
#   xlabel='Date',
#   ylabel='Count',
#   kind='line'
# )

```

```
from plotnine import aes, facet_wrap, ggplot, scale_x_datetime, options, stat_smooth, geom_col
```

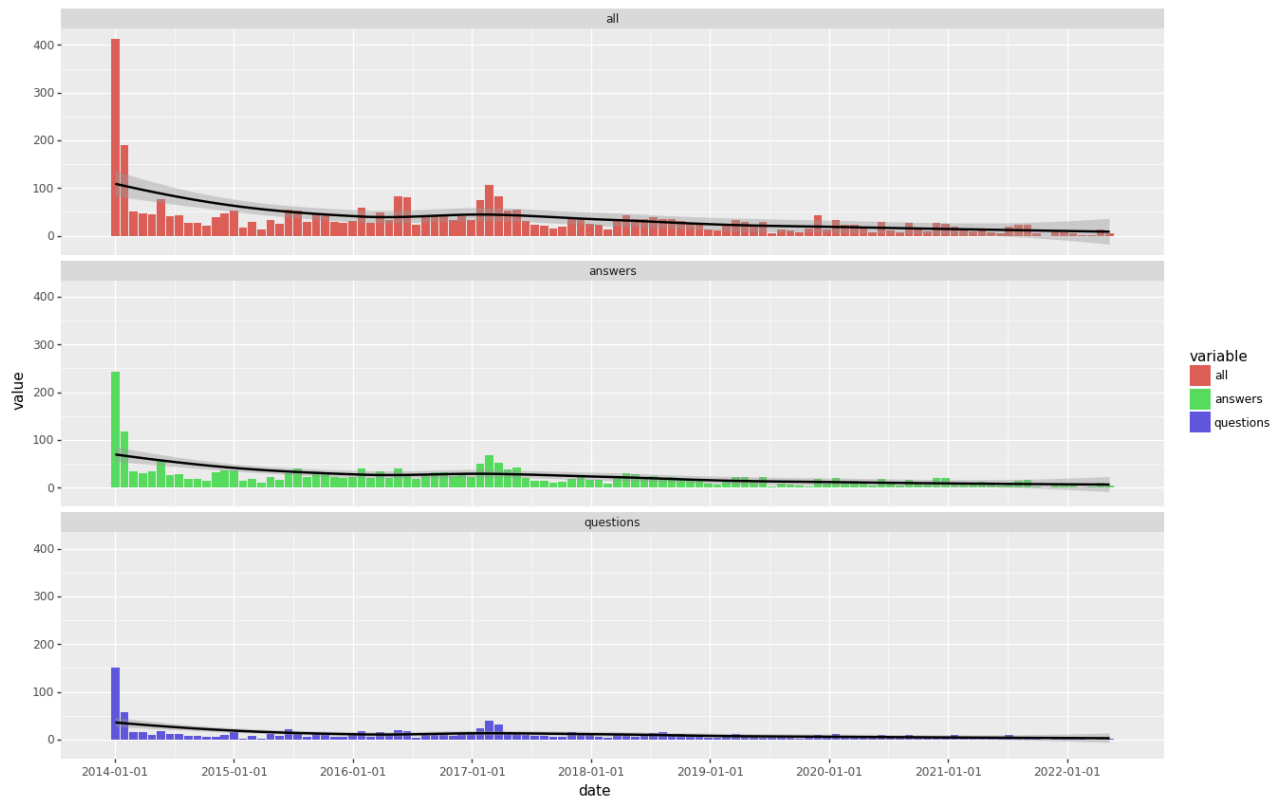
```
options.figure_size = (15, 10)
```

```
posts_long = posts_grouped.melt(id_vars=('date'), value_vars=('all', 'questions', 'answers'))
posts_long.head()
```

```

(ggplot(posts_long, aes(x='date', y='value', group='variable'))
+ geom_col(aes(fill='variable'))
+ scale_x_datetime()
+ stat_smooth(method='loess')
+ facet_wrap('variable', ncol=1)
)

```



```
<ggplot: (8785994108971)>
```

3.2 Forum retention - time from account creation to last comments

```
users = spark.read.format('parquet').load("outputs/users").select(f.col('id'), f.col('creation_date'))
```

```
users.show()
```

id	creation_date	display_name
-1	2014-01-21 17:45:...	Community
1	2014-01-21 20:21:...	Geoff Dalgas
2	2014-01-21 20:22:...	Kasra Rahjerdi
3	2014-01-21 20:22:...	Adam Lear
4	2014-01-21 20:22:...	Arie Litovsky
5	2014-01-21 20:22:...	Brian Nickel
6	2014-01-21 20:23:...	Jeremy T
7	2014-01-21 20:24:...	Tom Medley
8	2014-01-21 20:25:...	LessPop_MoreFizz
9	2014-01-21 20:25:...	Nick Craver
10	2014-01-21 20:28:...	ChrisG
11	2014-01-21 20:28:...	hairboat
12	2014-01-21 20:29:...	nhaarman
13	2014-01-21 20:29:...	Shog9
14	2014-01-21 20:32:...	Ben Collins


```
| 15|2014-01-21 20:33:...|          Ana|
| 16|2014-01-21 20:34:...|      Grace Note|
| 17|2014-01-21 20:34:...|      Jon Ericson|
| 18|2014-01-21 20:36:...|      awesome|
| 19|2014-01-21 20:36:...|  Steve Robbins|
+-----+
```

only showing top 20 rows

```
posts_by_user = posts.select(f.col('owner_user_id'), f.col('last_activity_date'), f.col('id').alias('id'))
```

```
posts_and_users_joined = (users
    .filter(f.col('id') != -1) # remove bots
    .join(posts_by_user, users.id == posts_by_user.owner_user_id, how="left" )
    .filter(f.col('post_id').isNotNull()) # remove users that never posted
)
```

```
posts_and_users_joined.show()
```

```
+-----+
| id|      creation_date| display_name|owner_user_id| last_activity_date|post_id|
+-----+
| 2|2014-01-21 20:22:...|Kasra Rahjerdi|          2|2014-01-22 00:26:...|    39|
| 2|2014-01-21 20:22:...|Kasra Rahjerdi|          2|2014-01-22 05:50:...|    28|
| 4|2014-01-21 20:22:...|Arie Litovsky|          4|2016-03-06 04:31:...|    85|
| 5|2014-01-21 20:22:...|Brian Nickel|          5|2014-02-04 23:20:...|   533|
| 5|2014-01-21 20:22:...|Brian Nickel|          5|2014-01-23 16:41:...|   267|
| 5|2014-01-21 20:22:...|Brian Nickel|          5|2014-01-22 18:36:...|   217|
| 5|2014-01-21 20:22:...|Brian Nickel|          5|2018-08-09 15:38:...|    50|
| 5|2014-01-21 20:22:...|Brian Nickel|          5|2014-01-21 20:59:...|    32|
| 5|2014-01-21 20:22:...|Brian Nickel|          5|2014-01-21 20:45:...|    17|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2014-01-29 20:34:...|   426|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2014-11-19 15:11:...|    82|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2014-01-21 21:47:...|    70|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2020-08-28 07:35:...|    59|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2022-01-14 10:04:...|    38|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2014-01-22 06:24:...|    35|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2014-01-22 17:04:...|    10|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2017-08-24 06:53:...|     8|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2021-01-15 06:17:...|     7|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2017-06-07 11:10:...|     5|
| 7|2014-01-21 20:24:...|Tom Medley|          7|2015-01-29 14:50:...|     3|
+-----+
```

only showing top 20 rows

```
posts_and_users_joined.select(f.col('post_id')).count() == posts_and_users_joined.select(f.col('post_id')).count()
```

True

```
user_last_post = (posts_and_users_joined
    .groupBy(f.col('id'), f.col('creation_date'))
    .agg(f.col('last_activity_date').max().alias('last_post')))
```

```

    .agg(
      f.max(f.col('last_activity_date'))
    )
  )

# time from account creation to last activity
user_last_post = user_last_post.withColumn('diff',f.datediff(f.col('max(last_activity_date)'), f.col('creation_date')))
user_last_post.show()

```

```

+-----+-----+-----+-----+
|  id|      creation_date|max(last_activity_date)|diff|
+-----+-----+-----+-----+
| 6696|2017-04-27 18:46:...| 2017-09-18 21:40:...| 144|
| 7212|2017-10-24 01:20:...| 2017-10-24 01:20:...|  0|
| 7311|2017-11-28 23:29:...| 2017-12-23 14:32:...| 25|
|10039|2020-01-17 20:47:...| 2020-03-25 21:19:...| 68|
|  149|2014-01-22 16:41:...| 2014-01-23 08:52:...|  1|
|  736|2014-04-03 13:25:...| 2017-01-27 14:13:...|1030|
| 4197|2015-06-10 19:13:...| 2015-06-10 19:18:...|  0|
| 5654|2016-07-13 09:17:...| 2016-09-02 12:08:...| 51|
| 7154|2017-10-01 06:24:...| 2017-10-01 06:24:...|  0|
| 7286|2017-11-23 11:43:...| 2017-11-24 22:07:...|  1|
| 7936|2018-07-20 09:31:...| 2018-07-24 11:49:...|  4|
|11698|2020-11-12 20:24:...| 2020-11-12 20:24:...|  0|
| 1295|2014-09-10 17:56:...| 2016-10-08 14:17:...| 759|
| 5893|2016-09-11 03:06:...| 2016-09-11 03:07:...|  0|
| 6636|2017-04-06 13:23:...| 2017-04-06 13:23:...|  0|
| 6699|2017-04-28 07:43:...| 2017-11-16 18:08:...| 202|
| 7208|2017-10-22 23:11:...| 2017-10-22 23:48:...|  0|
| 8088|2018-09-15 08:09:...| 2018-09-15 08:19:...|  0|
|  740|2014-04-04 15:58:...| 2016-06-21 13:46:...| 809|
| 1077|2014-07-09 23:08:...| 2018-10-09 20:22:...|1553|
+-----+-----+-----+-----+

```

only showing top 20 rows

```

user_last_post.select(f.col('id')).count() == user_last_post.select(f.col('id')).distinct().count()

```

True

```

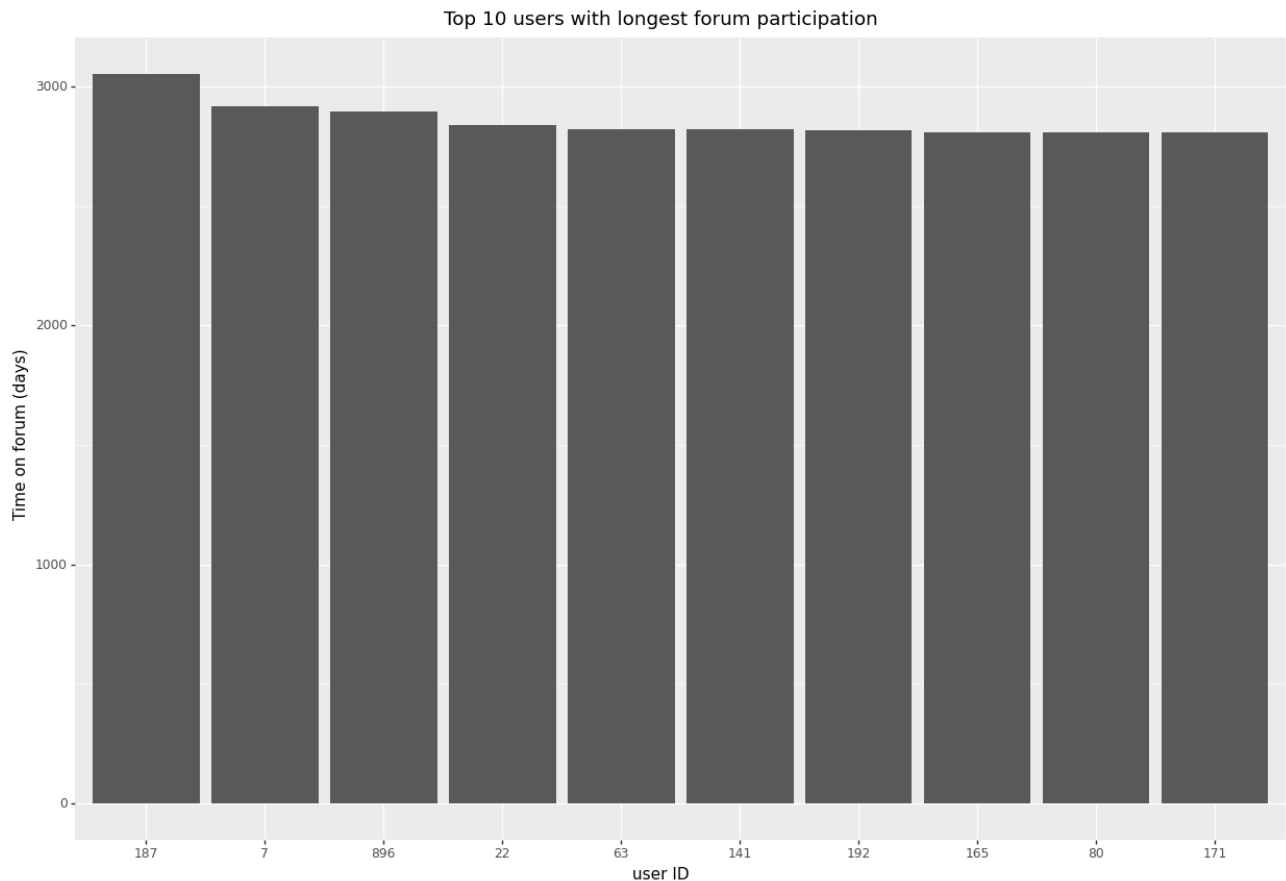
user_last_post_df = user_last_post.orderBy(f.col('diff').desc()).limit(10).withColumn('id_cat', f.col('id').cast('string'))

import pandas as pd
# add sorted categories for pretty plotting
user_last_post_df['id_cat'] = pd.Categorical(user_last_post_df.id_cat, categories=user_last_post_df.id_cat.sort_values().unique())

from plotnine import labs
(ggplot(user_last_post_df, aes(x='id_cat', y='diff'))
 + geom_col()
 + labs(x='user ID', y='Time on forum (days)', title='Top 10 users with longest forum participation')
)

```

/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: FutureWarning
/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: FutureWarning



<ggplot: (8785942620109)>

3.3 porównanie najwyżej i najniżej ocenianych pytań (długość, tagi, liczba odpowiedzi)

```
#users2 = spark.read.format('parquet').load("outputs/users").select(f.col('id'), f.col('creation_date'))  
#posts.show(1, vertical=True)
```

```
#1 - Question 2 - Answer 3 - Wiki 4 - TagWikiExcerpt 5 - TagWiki 6 - ModeratorNomination 7 - WikiPla  
questions = posts.select(f.col('id'), f.col('body_clean'), f.col('answer_count'), f.col('view_count'))  
                .filter(f.col('post_type_id') == 1)\  
                .drop(f.col('post_type_id'))
```


```
n_questions = 100  
top_questions = questions.orderBy(f.col('score'), ascending=False).limit(n_questions).withColumn('type', lit(1))  
bottom_questions = questions.orderBy(f.col('score'), ascending=True).limit(n_questions).withColumn('type', lit(2))
```

```
edge_questions = top_questions.unionAll(bottom_questions)
```

```
from pyspark.sql.functions import length
edge_questions = edge_questions.withColumn('post_len', f.length(f.col('body_clean')))
edge_questions_pd = edge_questions.toPandas()
```

```
from plotnine import ggplot, aes, geom_point, ggtitle
```

```
(ggplot(edge_questions_pd, aes(x = 'post_len', y = 'score')) \
+ geom_point() \
+ ggtitle(f'Top and bottom {n_questions} questions in relation to its length'))
```



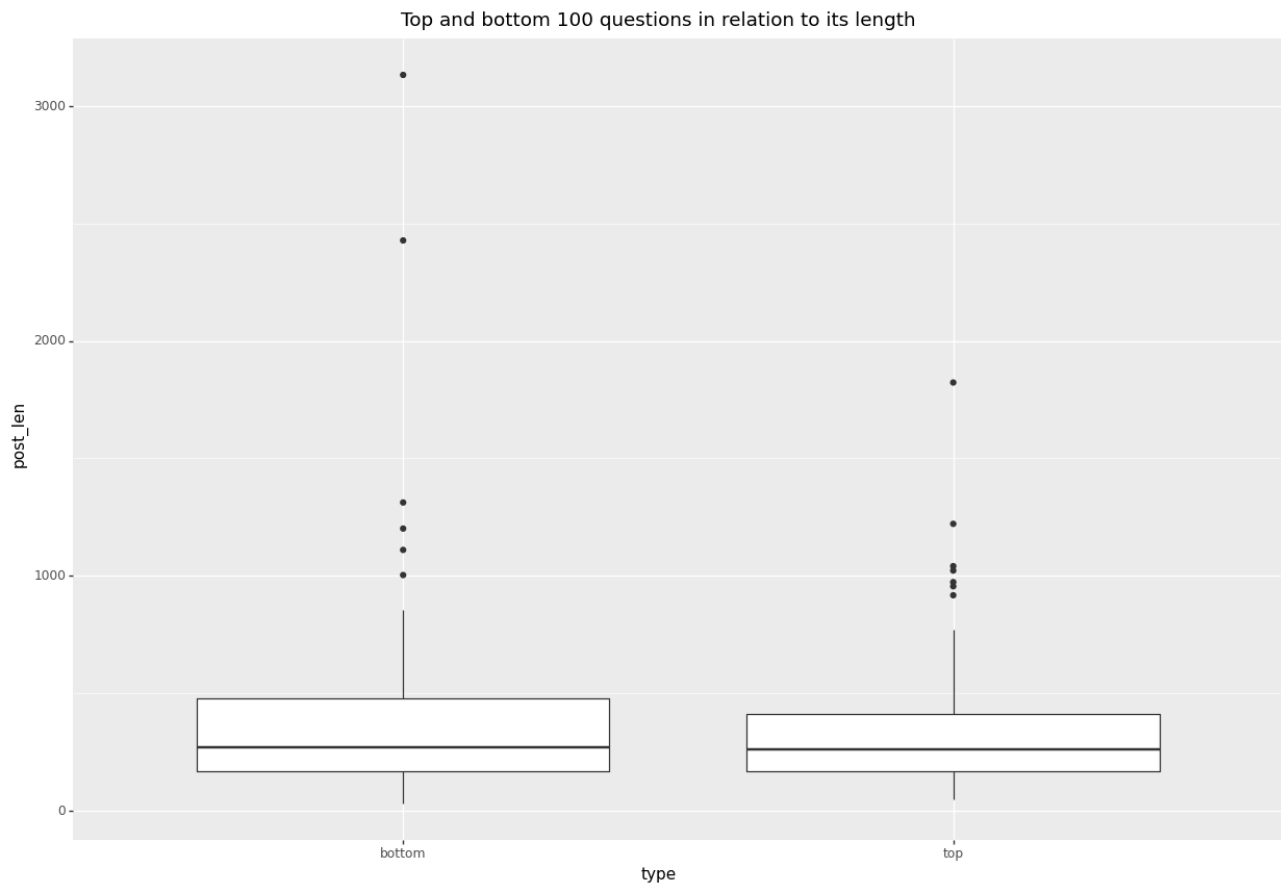
```
<ggplot: (8786045558433)>
```

```
edge_questions.groupby('type')\
    .agg(\
        f.max(f.col('post_len')),
        f.min(f.col('post_len')),
        f.mean(f.col('post_len')),
        f.stddev(f.col('post_len')),
        f.percentile_approx(f.col('post_len'), 0.5)
    ).show()
```

	type	max(post_len)	min(post_len)	avg(post_len)	stddev_samp(post_len)	percentile_approx(post_len, 0.5)
top		1823	46	349.08	275.08858738430234	
bottom		3133	30	389.26	424.65483595146975	

```
from plotnine import ggplot, aes, geom_boxplot, ggtitle


(ggplot(edge_questions_pd, aes(x = 'type', y = 'post_len')) \
 + geom_boxplot() \
 + ggtitle(f'Top and bottom {n_questions} questions in relation to its length'))
```



<ggplot: (8785942555621)>

```
from plotnine import ggplot, aes, geom_point, ggtitle, scale_x_log10, geom_smooth

(ggplot(edge_questions_pd, aes(x = 'view_count', y = 'answer_count', fill = 'type', size = 'score')) \
 + geom_point() \
 + scale_x_log10() \
 + geom_smooth() \
 + ggtitle(f'Top and bottom {n_questions} questions in relation to its views and answers'))
```



```
<ggplot: (8785942667757)>
```

```
edge_questions.groupby('type')\
    .agg(
        f.max(f.col('answer_count')),
        f.min(f.col('answer_count')),
        f.mean(f.col('answer_count')),
        f.stddev(f.col('answer_count')),
        f.percentile_approx(f.col('answer_count'), 0.5)
    ).show()
```

type	max(answer_count)	min(answer_count)	avg(answer_count)	stddev_samp(answer_count)	percentile_approx(answer_count, 0.5)
top	17	1	3.96	2.581675910150324	
bottom	4	0	1.28	0.9648363026488436	

```
edge_questions.groupby('type')\
    .agg(
        f.max(f.col('view_count')),
        f.min(f.col('view_count')),
        f.mean(f.col('view_count')),
        f.stddev(f.col('view_count')),
        f.percentile_approx(f.col('view_count'), 0.5)
    ).show()
```

	type	max(view_count)	min(view_count)	avg(view_count)	stddev_samp(view_count)	percentile_approx(view_count)
top		648941	175	26047.65	76276.20019090576	
bottom		9124	3	495.56	1220.7562924910794	

#<https://gist.github.com/dannymeijer/be3534470b205280e52dbbcbb19a9670>

```
from pyspark.sql import DataFrame
from pyspark.sql import functions as f
```

```
def regexp_extract_all(
    df: DataFrame,
    regex: str,
    no_of_extracts: int,
    input_column_name: str,
    output_column_name: str = "output",
    empty_array_replace: bool = True,
):
    """Pyspark implementation for extracting all matches of a reg_exp_extract
```

Background

The regular implementation of `regexp_extract` (as part of `pyspark.sql.functions` module) is not capable of returning more than 1 match on a `regexp` string at a time. This function can be used to circumvent this limitation.

How it works

You can specify a `no_of_extracts` which will essentially run the `regexp_extract` function that number of times on the `input_column` of the `df` (`DataFrame`). In between extracts, a set of interim columns are created where every intermediate match is stored. A distinct array is created from these matches, after which the interim columns are dropped. The resulting array is stored in the defined `output_column`. Empty strings/values in the resulting array can optionally be dropped or kept depending on how `empty_array_replace` is set (default is `True`).

Usage example

In the below example, we are extracting all email-addresses from a body of text. The returned `DataFrame` will have a new `ArrayType` column added named `email_addresses`

```
> # Assuming `df` is a valid DataFrame containing a column named `text`
> email_regex = r"[\w.-]+@[ \w.-]+\.[a-zA-Z]{1,}"
> df = regexp_extract_all(df, email_regex, 6, "text", "email_addresses", True)
```

Parameters

```
df: DataFrame
    Input DataFrame
```

```

regex: str
    Regexp string to extract from input DataFrame

no_of_extracts: int
    Max number of occurrences to extract

input_column_name: str
    Name of the input column

output_column_name: str
    Name of the output column (default: output)

empty_array_replace: bool
    If set to True, will replace empty arrays with null values (default: True)
"""
repeats = range(0, no_of_extracts)

# A set of interim columns are created that will be dropped afterwards
match_columns = [f"__{r}__" for r in repeats]

# Apply regexp_extract an r number of times
for r in repeats:
    df = df.withColumn(
        match_columns[r],
        f.regexp_extract(
            f.col(input_column_name),
            # the input regex string is amended with ".*?"
            # and repeated an r number of times
            # r needs to be +1 as matching groups are 1-indexed
            f".join([f'{regex}.*?' for i in range(0, r + 1)]),
            r + 1,
        ),
    )

# Create a distinct array with all empty strings removed
df = df.withColumn(
    output_column_name,
    f.array_remove(f.array_distinct(f.array(match_columns)), ""),
)

# Replace empty string with None if empty_array_replace was set
if empty_array_replace:
    df = df.withColumn(
        output_column_name,
        f.when(f.size(output_column_name) == 0, f.lit(None)).otherwise(
            f.col(output_column_name)
        ),
    )

# Drop interim columns
for c in match_columns:
    df = df.drop(c)

return df

```



```
#edge_questions.select(f.col('tags')).withColumn('tags_split', f.regexp_extract(f.col('tags'), r'<(\w+)>', 1))

edge_questions = regexp_extract_all(edge_questions, r'<(\w+)>', 99, "tags", "tags_split", True)

import pyspark.rdd as rdd
h = edge_questions.filter(f.col('type') == 'top').select(f.col('tags_split')).rdd
l = edge_questions.filter(f.col('type') == 'bottom').select(f.col('tags_split')).rdd

h.flatMap(lambda x: [y if y is not None else "" for y in x])\
  .flatMap(lambda x: [x[y] for y in range(0, len(x))])\
  .map(lambda x: (x, 1))\
  .aggregateByKey(0, (lambda acc,x: acc + x ), (lambda acc1,acc2: acc1+acc2))\
  .filter(lambda x: x[1] > 1)\
  .sortBy(lambda x: x[1], ascending=False)\
  .collect()
```

```
[('taste', 17),
 ('brewing', 14),
 ('history', 12),
 ('glassware', 8),
 ('storage', 8),
 ('serving', 8),
 ('style', 8),
 ('temperature', 6),
 ('stout', 5),
 ('terminology', 5),
 ('aging', 4),
 ('health', 4),
 ('bottles', 4),
 ('ingredients', 4),
 ('breweries', 3),
 ('ipa', 3),
 ('classification', 3),
 ('whiskey', 3),
 ('tripel', 3),
 ('drinking', 3),
 ('bottling', 3),
 ('flavor', 3),
 ('colour', 3),
 ('aroma', 2),
 ('freshness', 2),
 ('ale', 2),
 ('lager', 2),
 ('preservation', 2),
 ('foam', 2),
 ('dubbel', 2),
 ('skunking', 2),
 ('laws', 2),
 ('draught', 2),
 ('pouring', 2),
 ('pairing', 2),
```

```
('keg', 2),  
( 'water', 2),  
( 'trappist', 2),  
( 'carbonation', 2)]
```

```
tags_rdd = h.flatMap(lambda x: [y if y is not None else "" for y in x])\
    .flatMap(lambda x: [x[y] for y in range(0, len(x))])
```

```
x = tags_rdd.collect()
tags_str = ''
for y in range(len(x)):
    tags_str += f"{x[y]} "
```

```
from wordcloud import WordCloud
import matplotlib.pyplot as plt
wc = WordCloud(background_color = 'white').generate(tags_str)
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wc)
plt.axis("off")
plt.tight_layout(pad = 0)
plt.show()
```



```
tags_rdd = 1.flatMap(lambda x: [y if y is not None else "" for y in x])\
    .flatMap(lambda x: [x[y] for y in range(0, len(x))])
```


```
x = tags_rdd.collect()
tags_str = ''
for y in range(len(x)):
    tags_str += f"{x[y]} "
```

```
from wordcloud import WordCloud
import matplotlib.pyplot as plt
```

```

wc = WordCloud(background_color='white').generate(tags_str)
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wc)
plt.axis("off")
plt.tight_layout(pad = 0)
plt.show()

```



```

1.flatMap(lambda x: [y if y is not None else "" for y in x])\
  .flatMap(lambda x: [x[y] for y in range(0, len(x))])\
  .map(lambda x: (x, 1))\
  .aggregateByKey(0, (lambda acc,x: acc + x ), (lambda acc1,acc2: acc1+acc2))\
  .filter(lambda x: x[1] > 1)\
  .sortBy(lambda x: x[1], ascending=False)\
  .collect()

```

```

[('wine', 20),
 ('recommendations', 14),
 ('health', 11),
 ('taste', 8),
 ('history', 8),
 ('spirits', 7),
 ('breweries', 6),
 ('storage', 5),
 ('drinking', 4),
 ('flavor', 4),
 ('brewing', 4),
 ('alcohol', 4),
 ('pairing', 3),
 ('science', 3),
 ('champagne', 3),
 ('distillation', 3),
 ('temperature', 2),

```

```
( 'glassware', 2),
( 'scotch', 2),
( 'hangover', 2),
( 'water', 2),
( 'rum', 2),
( 'liquor', 2),
( 'draught', 2),
( 'vodka', 2),
( 'mead', 2),
( 'carbonation', 2),
( 'canada', 2),
( 'drink', 2),
( 'recipes', 2)]
```

3.4 procent przypadków kiedy najwyżej oceniana odpowiedź to nie zaakceptowana odpowiedź

```
#1 - Question 2 - Answer 3 - Wiki 4 - TagWikiExcerpt 5 - TagWiki 6 - ModeratorNomination 7 - WikiPla
posts_tmp = posts.select(f.col('id'), f.col("parent_id"), f.col('accepted_answer_id'), f.col('answer_
questions = posts_tmp.filter('post_type_id == 1 and answer_count > 0')\
    .select(f.col('id').alias('q_id'), f.col('accepted_answer_id'))
```

```
answers = posts_tmp.filter(f.col('post_type_id') == 2)\
    .select(f.col('id').alias('a_id'), f.col('parent_id'), f.col('score'))
```

```
from pyspark.sql import Window
```

```
window_partition_agg = Window.partitionBy("q_id")
```

```
questions.join(answers, on=questions.q_id == answers.parent_id)\
    .sort(['q_id', 'a_id'])\
    .withColumn("max_score", f.max(f.col("score")).over(window_partition_agg))\
    .filter(f.col("score") == f.col("max_score"))\
    .filter(f.col("accepted_answer_id").isNotNull())\
    .withColumn("is_accepted_best", f.col("accepted_answer_id") == f.col("a_id"))\
    .agg(
        f.sum(f.col("is_accepted_best").cast("integer")).alias("sum"),
        f.count(f.col("q_id")).alias("count")
    )\
    .withColumn("percent", (f.col("count") - f.col("sum")) / f.col("count") * 100).show()
```

```
+---+-----+-----+
|sum|count|           percent|
+---+-----+-----+
|641|  735|12.789115646258503|
+---+-----+-----+
```

3.5 rozkład ocen odpowiedzi zaakceptowanych vs pozostałych (średnia, odchylenie, minimum, maksimum)


```
window_partition_agg = Window.partitionBy("q_id")

questions.join(answers, on=questions.q_id == answers.parent_id)\
  .sort(['q_id', 'a_id'])\
  .withColumn("is_accepted", f.col("accepted_answer_id") == f.col("a_id"))\
  .groupBy(f.col("is_accepted")).agg(
    f.avg(f.col("score")).alias("avg_score"),
    f.stddev(f.col("score")).alias("std_score"),
    f.min(f.col("score")).alias("min_score"),
    f.max(f.col("score")).alias("max_score"),
    f.count("a_id")
  ).show()
```

is_accepted	avg_score	std_score	min_score	max_score	count(a_id)
null	2.7551686615886832	3.1818372333580007	-5	30	919
true	6.395043731778426	5.915949387154137	0	46	686
false	2.5841694537346713	2.7353292123298076	-4	26	897

```
accepted_df = questions.join(answers, on=questions.q_id == answers.parent_id)\
  .sort(['q_id', 'a_id'])\
  .withColumn("is_accepted", (f.col("accepted_answer_id") == f.col("a_id")).cast("string"))\
  .withColumn("is_accepted", f.when(f.col("is_accepted").isNull(), "no accepted answer").otherwise(f.col("is_accepted")))

(ggplot(accepted_df, aes(x="is_accepted", y="score"))\
  +geom_boxplot()\
  +ggtitle("Distribution of accepted vs not accepted answers (including answers for questions witho"))
```



```
<ggplot: (8785942528877)>
```

3.6 top N tagów które wygenerowały najwięcej wyświetleń

```
tags_views = posts.select(['tags', 'view_count'])
tags_views_agg = regexp_extract_all(tags_views, r'<(\w+)>', 99, "tags", "tags_split", True)\
    .select([f.explode(f.col('tags_split')).alias("tag"), f.col("view_count")])\
    .filter(f.col("view_count").isNotNull())\
    .groupBy('tag')\
    .agg(
        f.sum("view_count").alias("sum_views")
    )
```

```
tag_top_views.head(20)
```

	tag	sum_views
0	taste	1330670
1	health	1286001
2	preservation	682216
3	storage	542860
4	whiskey	464756
5	bourbon	330268

	tag	sum_views
6	brewing	307892
7	ipa	291935
8	spirits	255328
9	drinking	225924
10	temperature	218203
11	drink	204991
12	tequila	196689
13	alcohol	188615
14	recommendations	185154
15	wine	181081
16	style	168681
17	flavor	168594
18	history	167414
19	pairing	164589

3.7 liczba postów w czasie dla każdego z top N tagów (lineplot/barplot)

```
#1 - Question 2 - Answer 3 - Wiki 4 - TagWikiExcerpt 5 - TagWiki 6 - ModeratorNomination 7 - WikiPla
posts_tmp = posts.select(f.col('id'), f.col('creation_date'), f.col('tags'))

posts_tags_time = regexp_extract_all(posts_tmp, r'<(\w+)>', 99, "tags", "tags_split", True).withColumn

top_posts_tags_time = posts_tags_time.join(tag_top_views, on="tag", how="inner").select(f.col('id'),

(ggplot(top_posts_tags_time_agg_pd, aes("date", "count"))\
  + scale_x_datetime()\
  + geom_col() \
  + facet_wrap("tag", ncol=3) \
  + ylim(0, 15))
```

3.8 najczęściej pojawiające się słowa w tytułach (z pominięciem stopwords)

```
from bs4 import BeautifulSoup
from html import unescape
from pyspark.sql.functions import udf, regexp_replace
from pyspark.sql.types import *
# remove html tags
def tags_remove(s):
    soup = BeautifulSoup(unescape(s), 'lxml')
    return soup.text

udf_tags_remove = udf(lambda m: tags_remove(m))

titles = posts.filter(f.col("title").isNotNull()).select(f.col("title"))\
    .withColumn("title_clean", f.lower(f.col("title")))\
    .withColumn("title_clean", regexp_replace('title_clean', "[^a-zA-Z\\s]", " "))
```

```

from pyspark.ml.feature import Tokenizer, StopWordsRemover
from nltk.stem.snowball import SnowballStemmer
udf_filter_length = udf(lambda row: [x for x in row if len(x) > 1], ArrayType(StringType()))

stemmer = SnowballStemmer(language='english')
stemmer_udf = udf(lambda token: stemmer.stem(token), StringType())

tokenizer = Tokenizer(inputCol='title_clean', outputCol='words_token')
title_tokens = tokenizer.transform(titles).withColumn('words_token', udf_filter_length(f.col('words_token')))

remover = StopWordsRemover(inputCol='words_token', outputCol='words_no_stop')
title_tokens_no_stop = remover.transform(title_tokens)
exploded = title_tokens_no_stop.withColumn("words", f.explode(f.col("words_no_stop")))

title_stem = exploded.withColumn('words_stem', stemmer_udf("words"))

word_lookup = title_stem.select([f.col("words"), f.col("words_stem")]).distinct()
word_lookup.show() # TODO aggregate this
title_stem.groupBy("words_stem").agg(f.count("title").alias('count')).orderBy('count', ascending=False)

```

```

+-----+-----+
|      words| words_stem|
+-----+-----+
|      opened|      open|
|antidepressants|antidepress|
|      taken|      taken|
|  alternative|  altern|
|    learning|    learn|
|    sherry|    sherri|
|   regionali| regionali|
|   archetype| archetyp|
|   inhibitor| inhibitor|
|    outside|    outsid|
|      bay|      bay|
|   sangria|   sangria|
|    invest|    invest|
|   togther|   togther|
|    fake|    fake|
|    kahlua|    kahlua|
|   imported|   import|
|    tables|    tabl|
|    desire|    desir|
|   bavaria|   bavaria|
+-----+-----+

```

only showing top 20 rows

```

+-----+-----+
|words_stem|count|

```



```

+-----+-----+
|      beer| 476|
|      wine| 147|
|     drink| 104|
| alcohol| 88|
|    differ| 72|
|    bottl| 68|
|      use| 50|
|     tast| 47|
|     brew| 43|
|     make| 41|
|     good| 33|
| cocktail| 29|
|      age| 27|
| recommend| 26|
|      ale| 26|
|     like| 24|
|     made| 23|
|  whiskey| 23|
|   spirit| 23|
|      one| 22|
+-----+-----+

```

only showing top 20 rows

3.9 procent użytkowników którzy nigdy nic nie zapostowali

```

# users.show(2)
# posts.show(2)

users_posts = (users.join(
  (posts.select(f.col('id').alias('post_id'), f.col('owner_user_id'))), on=[users.id == posts.owner
  .filter("id IS NOT NULL and NOT id = -1")
  )

users_posts.select([f.col('id'), f.col('post_id')])\
  .groupBy("id")\
  .agg(
    f.count(f.col('post_id')).alias('post_count')
  )\
  .agg(
    f.sum(f.when(f.col('post_count') == 0, f.lit(1)).otherwise(f.lit(0))).alias("not_posted"),
    f.count('id').alias('all')
  ) \
  .withColumn('% not posted', (f.col("not_posted") / f.col('all') * 100)).show()

# (users.join(
#   (posts.select(f.col('id').alias('post_id'), f.col('owner_user_id'))), on=[users.id == posts.own
#   .filter(f.col('id').isNotNull() & f.col('post_id').isNull())\
#   .select([f.col('id'), f.col('display_name')]).distinct()\
#   .show()
#   )

```

```

+-----+-----+

```

```
|not_posted| all|      % not posted|
+-----+-----+-----+
|      7691|8947|85.96177489661339|
+-----+-----+-----+
```

3.10 średni czas od pojawienia się pytania do pojawienia się zaakceptowanej odpowiedzi

```
# keep only questions with answers
questions = posts.filter(f.col('post_type_id') == 1).filter(f.col('answer_count') > 0).select([f.col('parent_id').alias('q_id'), f.col('creation_date').alias('q_creation_date')])
answers = posts.filter(f.col('post_type_id') == 2).select([f.col('id').alias('a_id'), f.col('parent_id').alias('q_id'), f.col('creation_date').alias('a_creation_date')])
#posts.show(1, vertical=True)
time_to_accept = questions.join(answers, on=[questions.accepted_answer_id==answers.a_id])\
    .withColumn('time_to_accept_sec', f.unix_timestamp('a_creation_date') - f.unix_timestamp('q_creation_date'))\
    .withColumn('time_to_accept_min', f.round(f.col('time_to_accept_sec') / 60, 2))


time_to_accept.agg(
    f.avg('time_to_accept_min'),
    f.stddev('time_to_accept_min'),
    f.percentile_approx("time_to_accept_min", [0.25, 0.5, 0.75], 1000000).alias("quantiles")
).show(truncate=False)

time_to_accept_pd = time_to_accept.withColumn('all', f.lit("all_questions")).toPandas()

from plotnine import geom_jitter
(ggplot(time_to_accept_pd, aes(x='all', y="time_to_accept_min"))\
 +geom_jitter())
```

```
+-----+-----+-----+
|avg(time_to_accept_min)|stddev_samp(time_to_accept_min)|quantiles|
+-----+-----+-----+
|25244.9435714286      |123338.6325101642      |[141.53, 753.25, 3605.8]|
+-----+-----+-----+
```

```
/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: FutureWarning:
/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: FutureWarning
```



<ggplot: (8785939115177)>

3.10.1 remove outliers


```
no_outliers = time_to_accept.filter(f.col('time_to_accept_min') < 100000)

no_outliers.agg(
  f.avg('time_to_accept_min'),
  f.stddev('time_to_accept_min'),
  f.percentile_approx("time_to_accept_min", [0.25, 0.5, 0.75], 1000000).alias("quantiles")
).show(truncate=False)

no_outliers_pd = no_outliers.withColumn('all', f.lit("all_questions")).toPandas()
(ggplot(no_outliers_pd, aes(x='all', y="time_to_accept_min"))\
  +geom_jitter())
```

avg(time_to_accept_min)	stddev_samp(time_to_accept_min)	quantiles
4769.308792048925	12694.627623679486	[127.72, 644.93, 2922.82]

/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: FutureWarning
/config/workspace/.venv/lib/python3.10/site-packages/pyspark/sql/pandas/conversion.py:248: FutureWarning



<ggplot: (8785939033767)>

References

4 Załączniki

4.1 Polecenia budujące infrastrukturę

4.1.1 EMR

```
aws emr create-cluster --name="MyEMRCluster" \  
  --release-label emr-6.8.0 \  
  --applications Name=JupyterHub Name=Hadoop Name=Spark \  
  --log-uri s3://emr-logs-beer-and-wine/MyJupyterClusterLogs \  
  --use-default-roles \  
  --instance-groups InstanceGroupType=MASTER,InstanceCount=1,InstanceType=m4.large InstanceGroupType=SLAVE,InstanceCount=2,InstanceType=m4.large \  
  --ebs-root-volume-size 32 \  
  --configurations file://emr-configurations.json \  
  --bootstrap-actions Path=s3://misc-beer-and-wine/install_python_libraries.sh,Name=InstallJupyterLibraries
```

Plik `install-my-jupyter-libraries.sh` dostępny jest pod poniższym adresem

4.1.2 S3

Koszyki S3 zostały utworzone przy pomocy poniższego polecenia:

```
aws s3api create-bucket --acl private --bucket <nazwa koszka>
```